

TCPS: A Task and Cache-Aware Partitioned Scheduler for Hard Real-Time Multi-core Systems

Yixian Shen, Jun Xiao and Andy D Pimentel


LCTES-2022

Outline

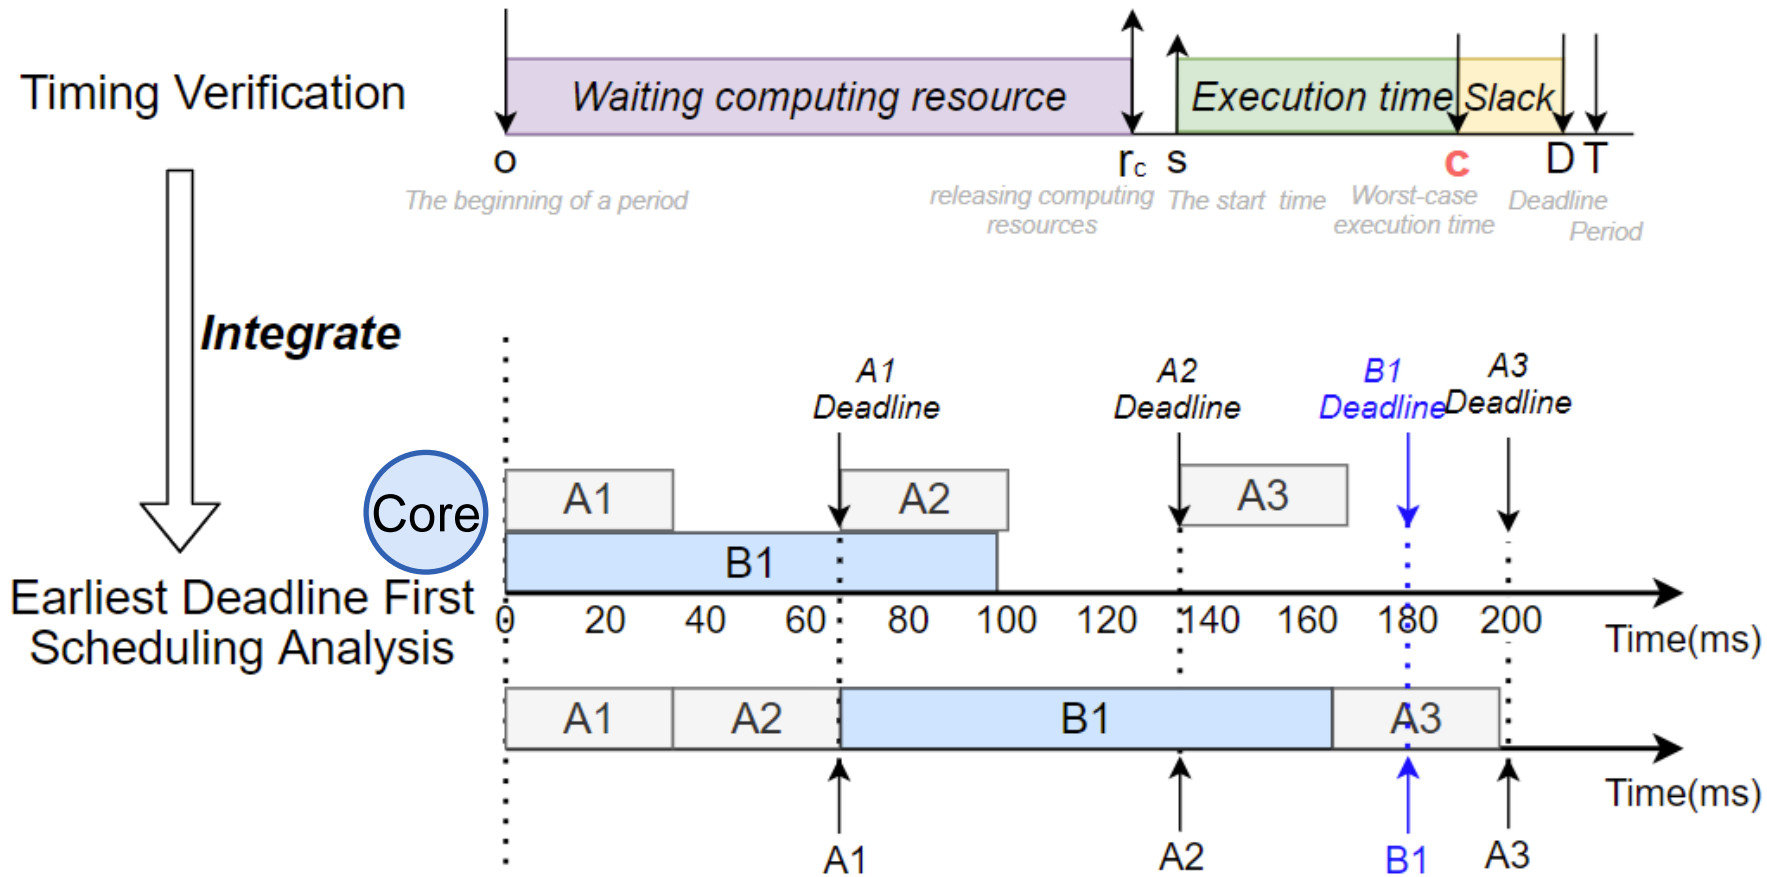
- **Research Background**
- **Start-of-art techniques**
- **Contribution**
- **TCPS approach**
- **Experiments**
- **Conclusion**

Background

- Time predictability is **crucial** for hard latency-critical applications
 - ❖ **Airbag restraint systems:** Car airbag, Electronic parachute
 - ❖ **Precise instrument:** Leonardo's Robot, Microscope
 - ❖ **Avionic systems:** Airbus, Boeing Co.
 - ❖ **Acute Sensing systems:** Earthquake sensor, Fire alarm sensor
 - ❖ **Unmanned aerial vehicles:** Drones

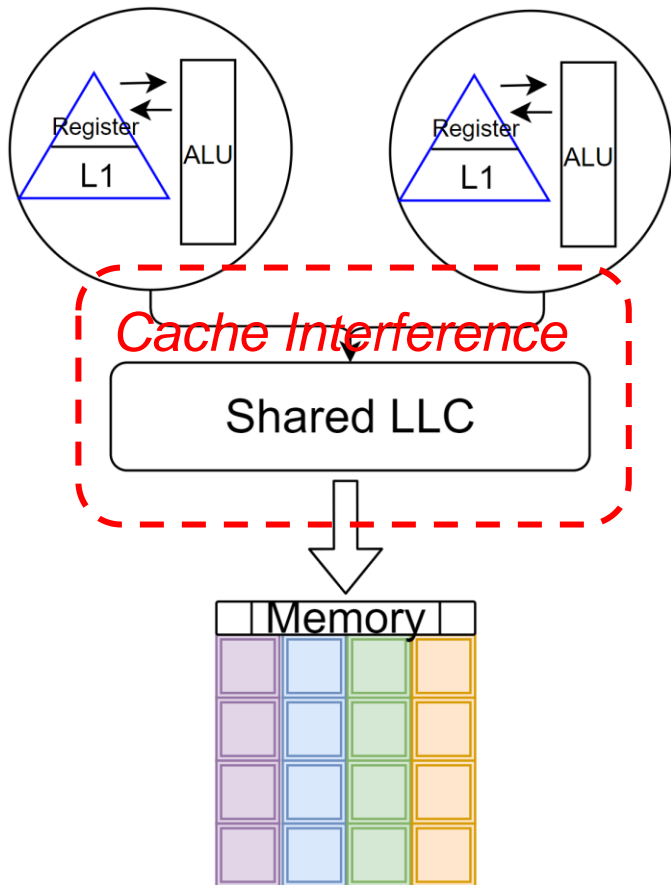


Timing verification in single-core real-time systems



Challenges

- ❑ Time accidents degrade time predictability



➤ Types of time accidents

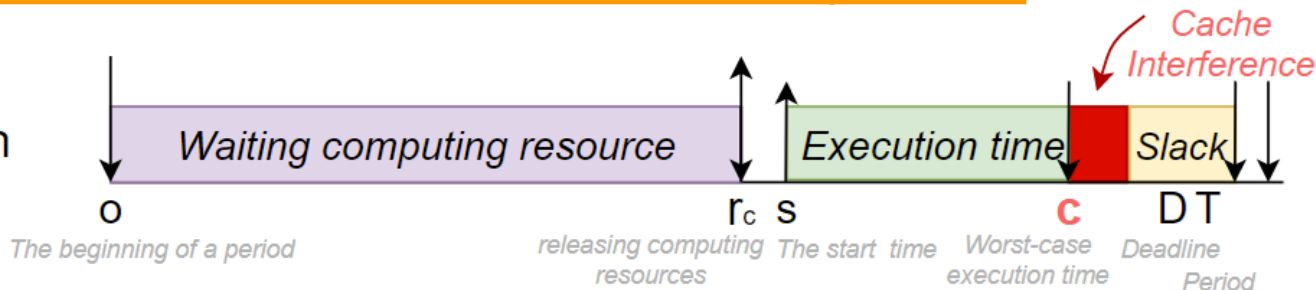
- ❖ Pipeline stalls
- ❖ Branch mispredictions
- ❖ Bus collisions
- ❖ Cache misses
- ❖ TLB misses
- ❖ Memory refresh of DRAM

Constant



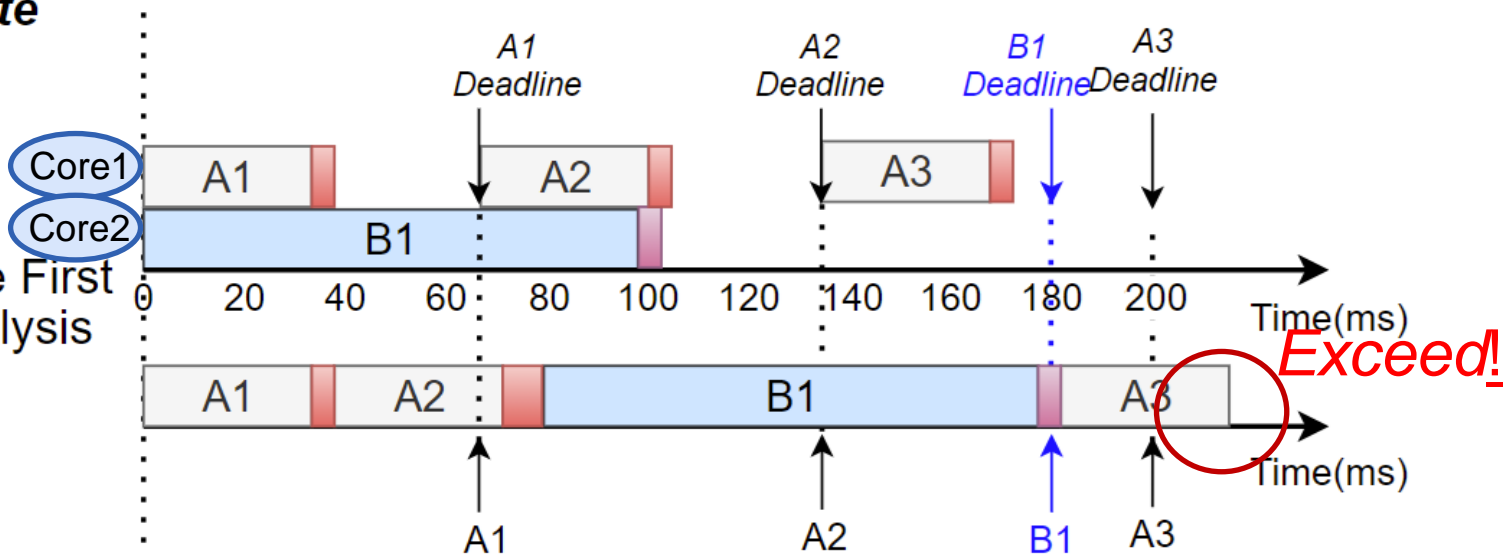
Timing verification in multi-core real-time systems

Timing Verification

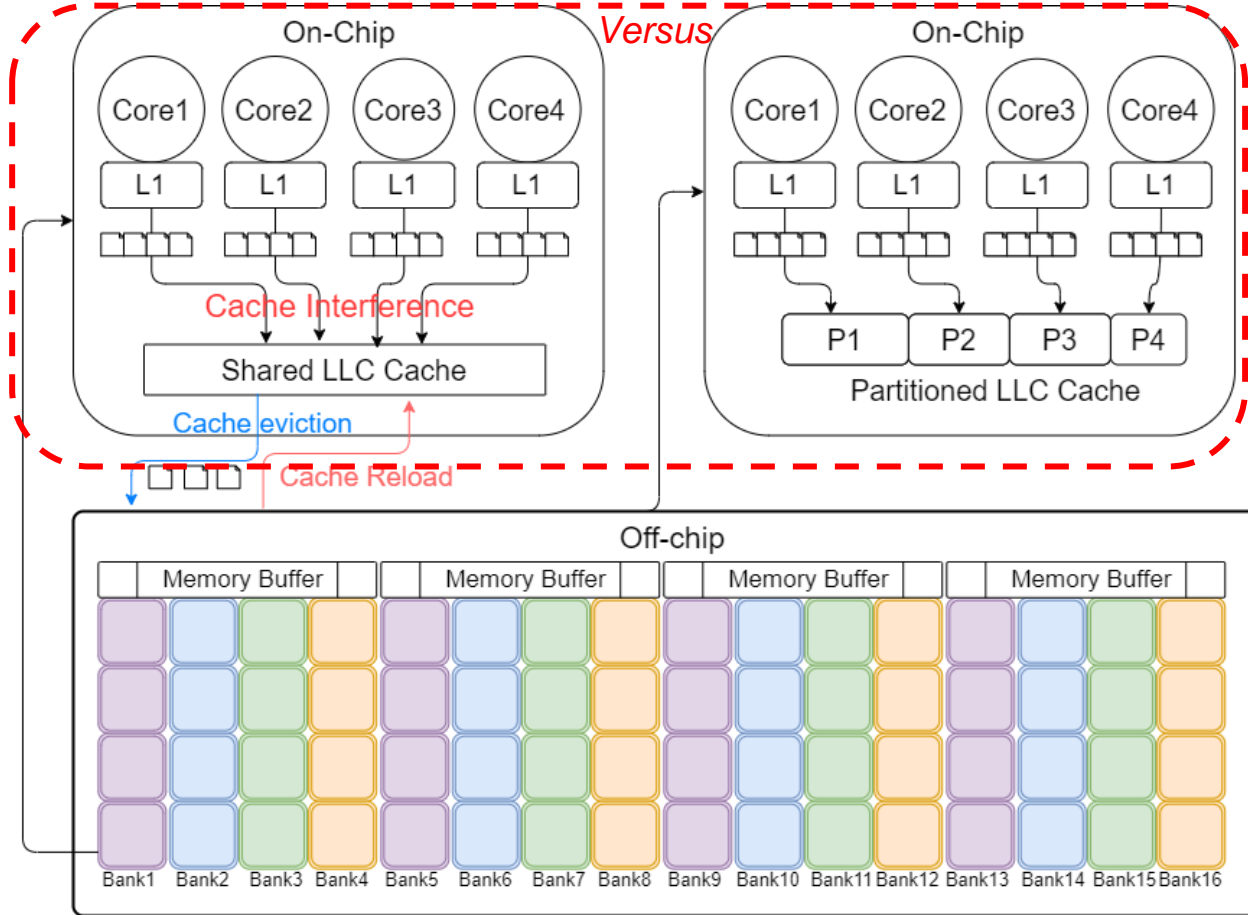


Integrate

Earliest Deadline First
Scheduling Analysis



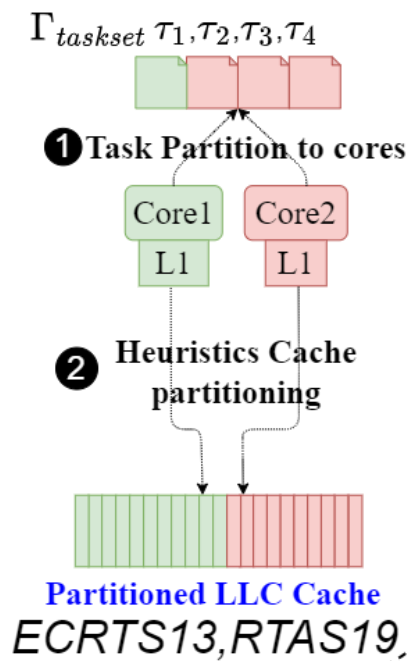

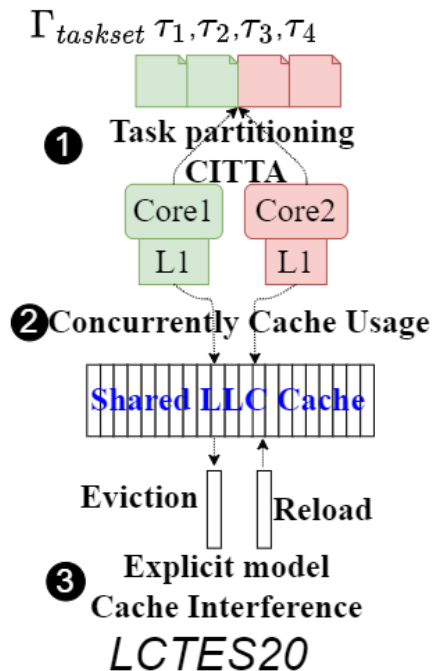
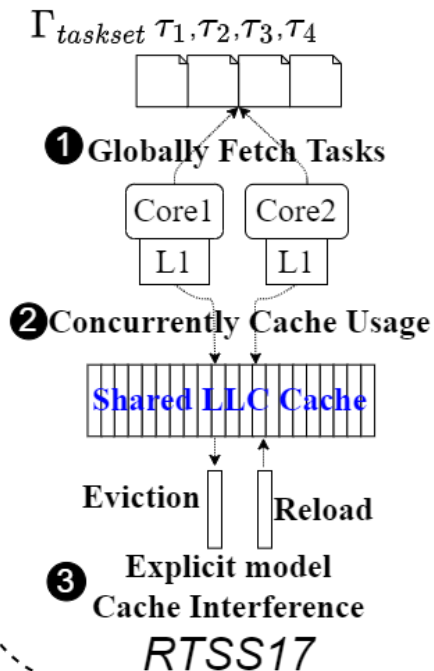
Two mainstream methods to improve the time predictability



- Explicitly calculate the cache inference
- Cache partitioning

Motivation

Start-of-art related works




Contribution

- TCPS: a new heuristic **task and cache-aware** partitioned scheduling policy for non-preemptive real-time multi-core systems, which combines the benefits of **cache partitioning and partitioned scheduling**.
- We evaluate the **design choices of our scheduling policy**, comparing it with all different combinations of methods to address cache interference and scheduling approaches.
- We conduct comprehensive experiments and **identify how different parameter settings** affect the relative performance of partitioned and non-partitioned shared caches for different real-time schedulers


System model

- Real-time task $\tau_k = (C_k, D_k, T_k)$
- Multiple real-time tasks $\tau = (\tau_1, \tau_2, \dots, \tau_n)$
- Multiple processor with m cores $\pi = (\pi_1, \pi_2, \dots, \pi_m)$
 - Shared caches
 - Cache interference
 - Inter-core cache interference
 - Intra-core cache interference
 - C_k does not count cache interference
- Scheduling policy
 - Non-preemptive EDF \rightarrow no intra-core cache interference

TCPS approach

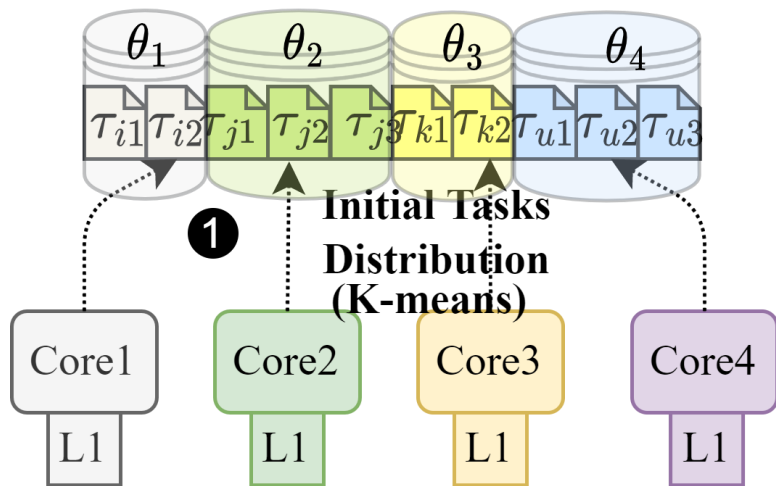


(a) Mälardalen benchmarks



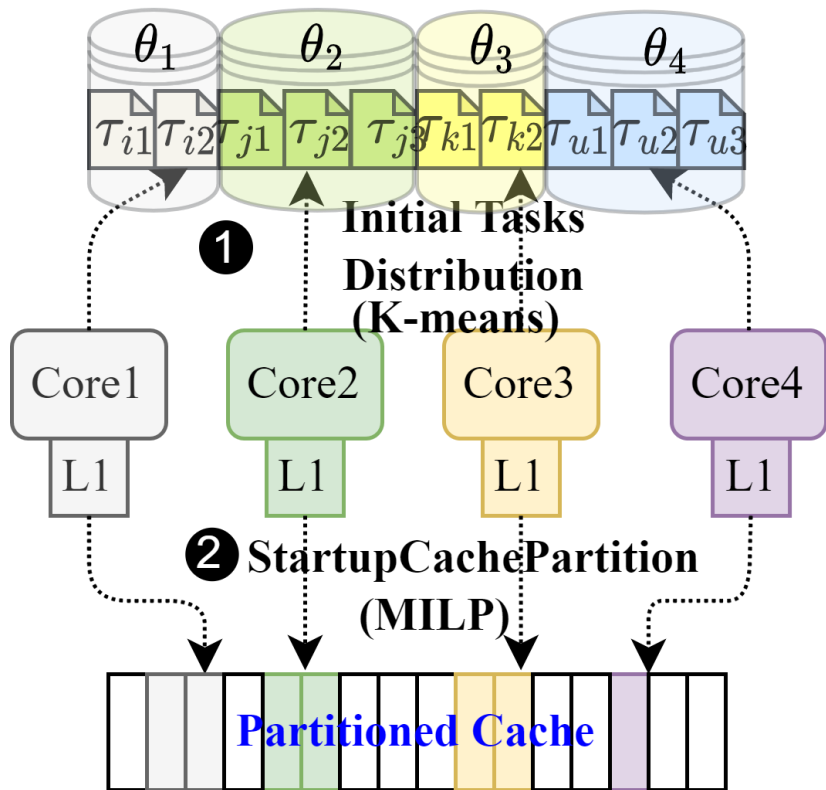
(b) TACLeBench

TCPS approach



$$\underset{\theta}{\operatorname{argmin}} \sum_{i=1}^m \frac{1}{2|\theta_i|} \sum_{\tau_k, \tau_j \in \theta_i} \|\vec{g}_k - \vec{g}_j\|^2$$

TCPS approach



$$\tau = \{\tau_1, \tau_2, \dots, \tau_n\}$$

$$\min\{U_1 + U_2 + \dots + U_n\}$$

$$\min\left\{\frac{1}{T_1} * c_1 + \frac{1}{T_2} * c_2, \dots, + \frac{1}{T_n} * c_n\right\}$$

Constraints:

$$p_1 + p_2 + \dots + p_n \leq n$$

$$c_1 * A_1 + p_1 * B_1 + k_1 = 0$$

$$c_2 * A_2 + p_2 * B_2 + k_2 = 0$$

...

$$c_n * A_n + p_n * B_n + k_n = 0$$

The cache partitioning algorithm

$$\tau = \{\tau_1, \tau_2, \dots, \tau_n\}$$

$$\min\{U_1 + U_2 + \dots + U_n\}$$
$$\min\{\frac{1}{T_1} * c_1 + \frac{1}{T_2} * c_2, \dots, + \frac{1}{T_n} * c_n\}$$

Constraints:

$$p_1 + p_2 + \dots + p_n \leq n$$

$$c_1 * A_1 + p_1 * B_1 + k_1 = 0$$

$$c_2 * A_2 + p_2 * B_2 + k_2 = 0$$

...

$$c_n * A_n + p_n * B_n + k_n = 0$$

$U_i \leftarrow$ utilization

$T_i \leftarrow$ Period

$c_i \leftarrow$ WCET

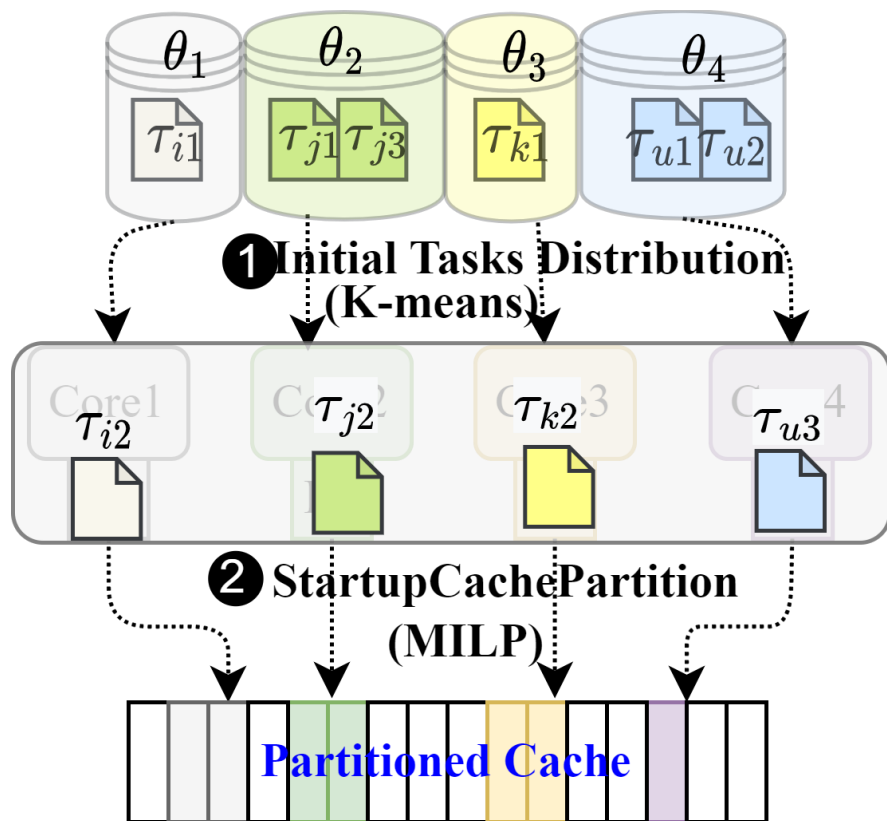
$$U_i = \frac{c_i}{T_i}$$

$p_i \leftarrow$ partition size

$A_i, B_i \leftarrow$ coefficient

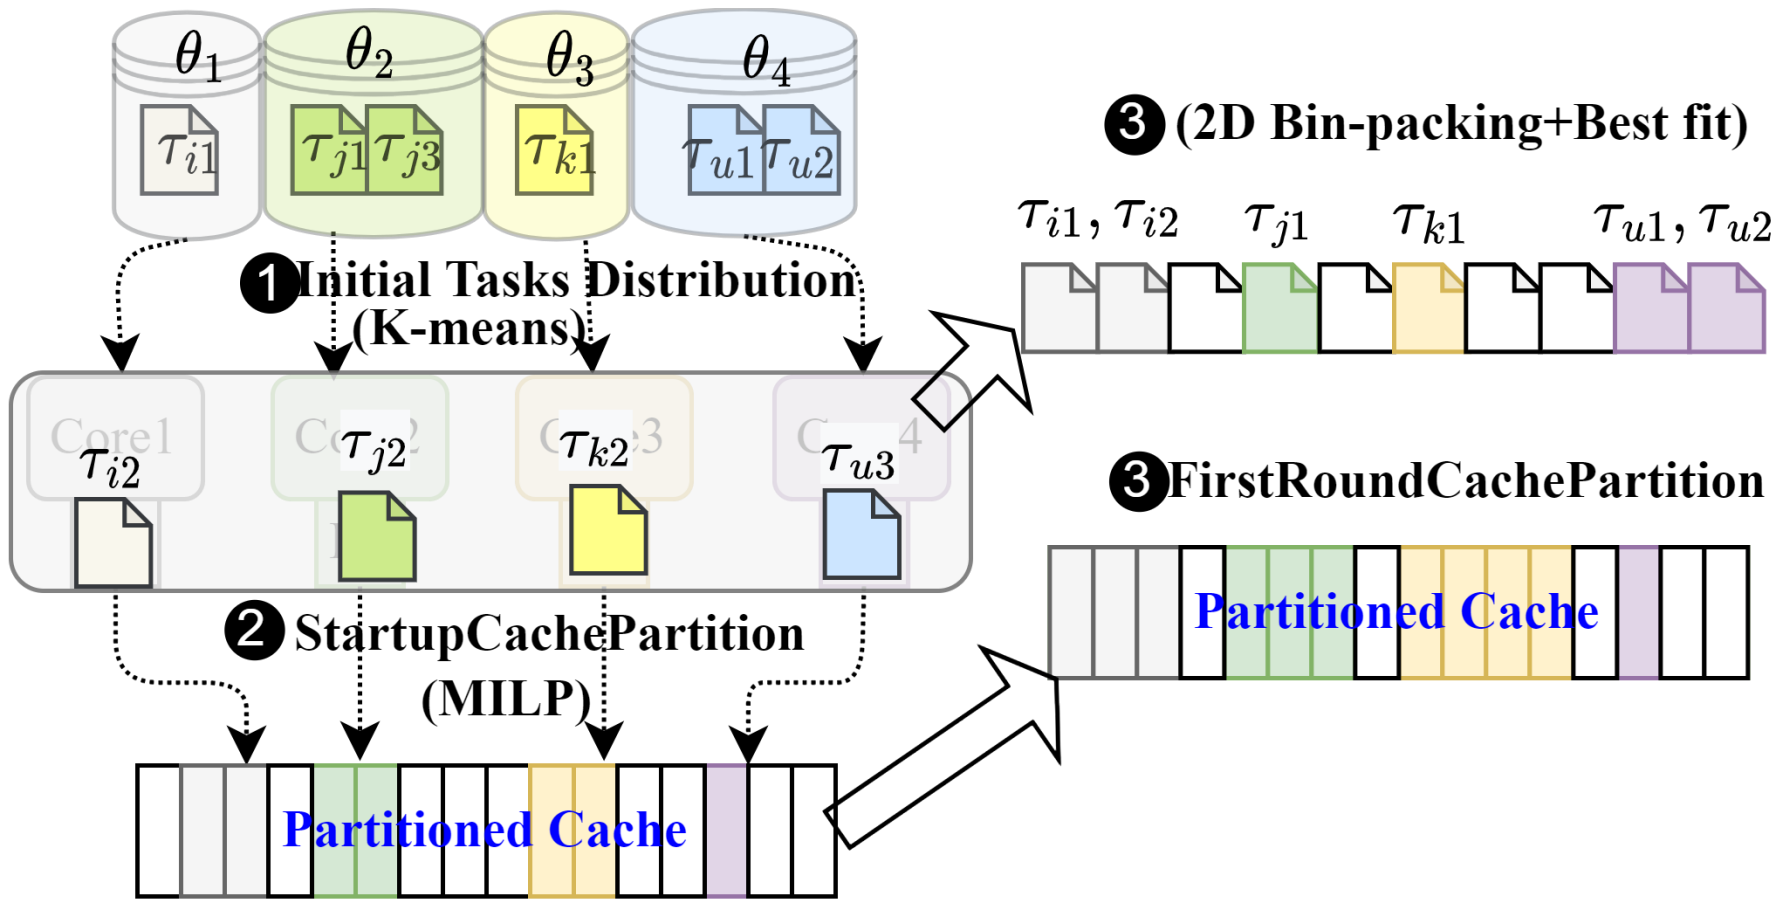
$k_i \leftarrow$ offset

TCPs approach

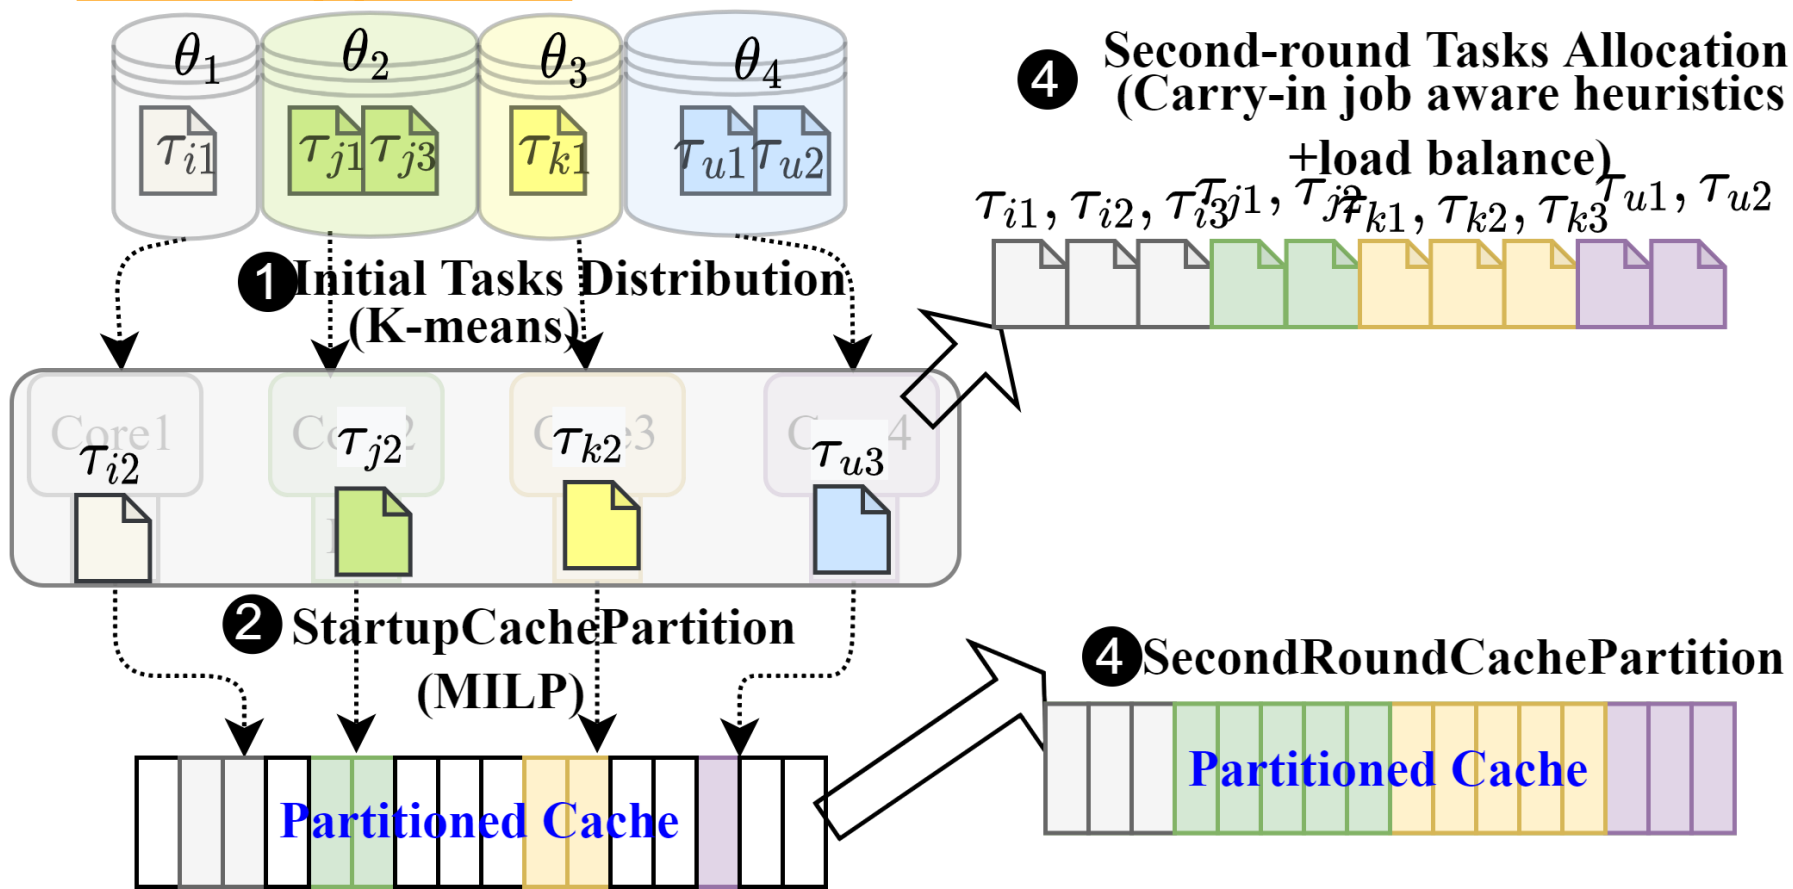


$$D_k \geq \sum_{\substack{\tau_j \in \tau(\pi_i) \cup \{\tau_k\} \\ D_j < D_k}} DBF^*(\tau_j, D_k) + \max_{\substack{\tau_j \in \tau(\pi_i) \cup \{\tau_k\} \\ D_j > D_k}} C_j$$

TCPS approach




TCPS approach




TCPS approach

Design time

3 First-round Task Allocation (2D Bin-packing+Best fit)

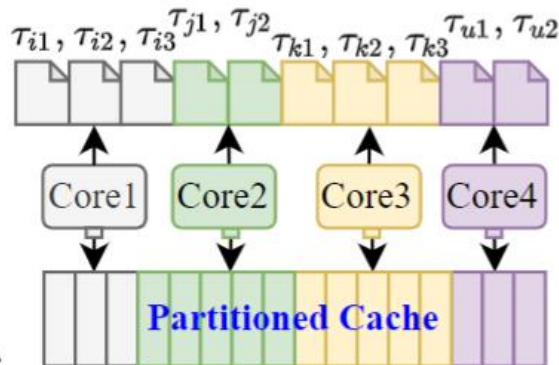


4 Second-round Tasks Allocation (Carry-in job aware heuristics +load balance)




Run time

5 Partitioned Scheduling




TaskBuckets

θ



1 Initial Tasks Distribution (K-means)



2 StartupCachePartition (MILP)



3 FirstRoundCachePartition



4 SecondRoundCachePartition




Experiments

- Workload generation policy
 - The number of tasks n
 - Total task utilization U
 - WCET: measured by Heptane
 - Cache interference: measured by the extended Heptane
- In each experiment, generating totally 20000 taskset
- Evaluation Metric
 - Acceptance ratio: the number of schedulable tasksets divided by the total number of tasksets.
 - Cache Usage
 - Load balance

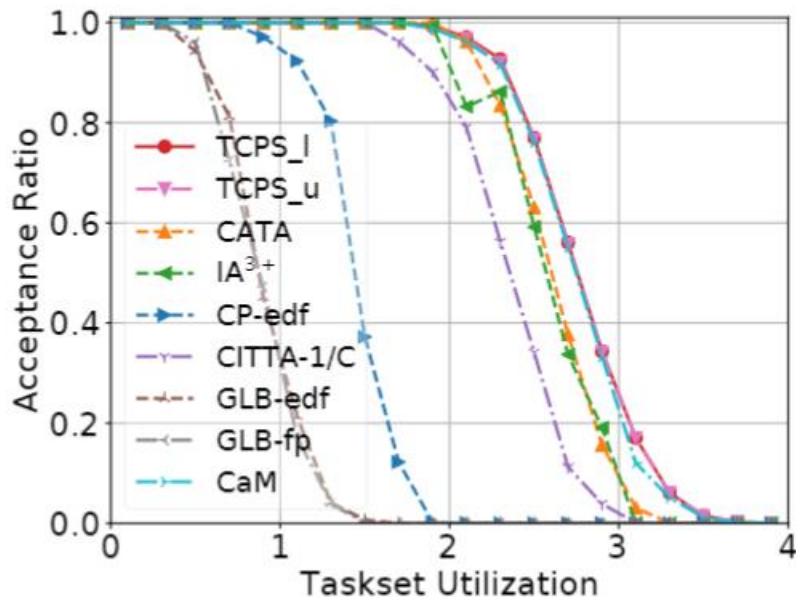
Experiments

Cache-sensitive tasks: their WCET drops by 40% or more when receiving 12.5% of cache space compared to the minimum allocation space.

cache-insensitive tasks: their WCET drops by 10% or less when receiving 50% cache space.



(a) $m=4, n=10$, cache-sensitive Γ



(b) $m=4, n=10$, cache-insensitive Γ

Experiments

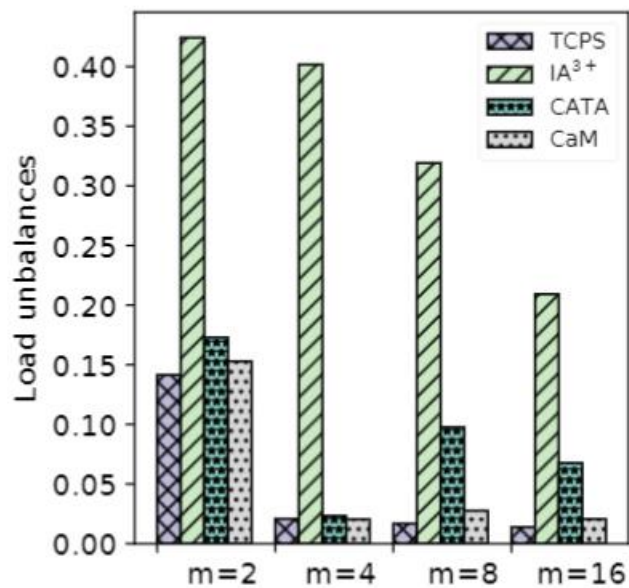
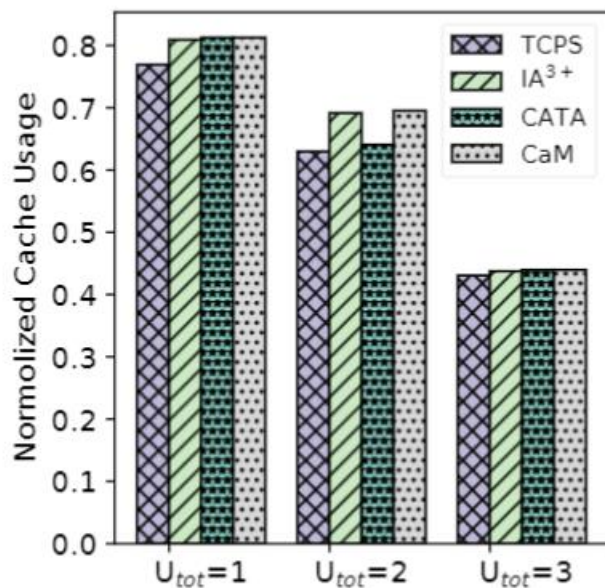
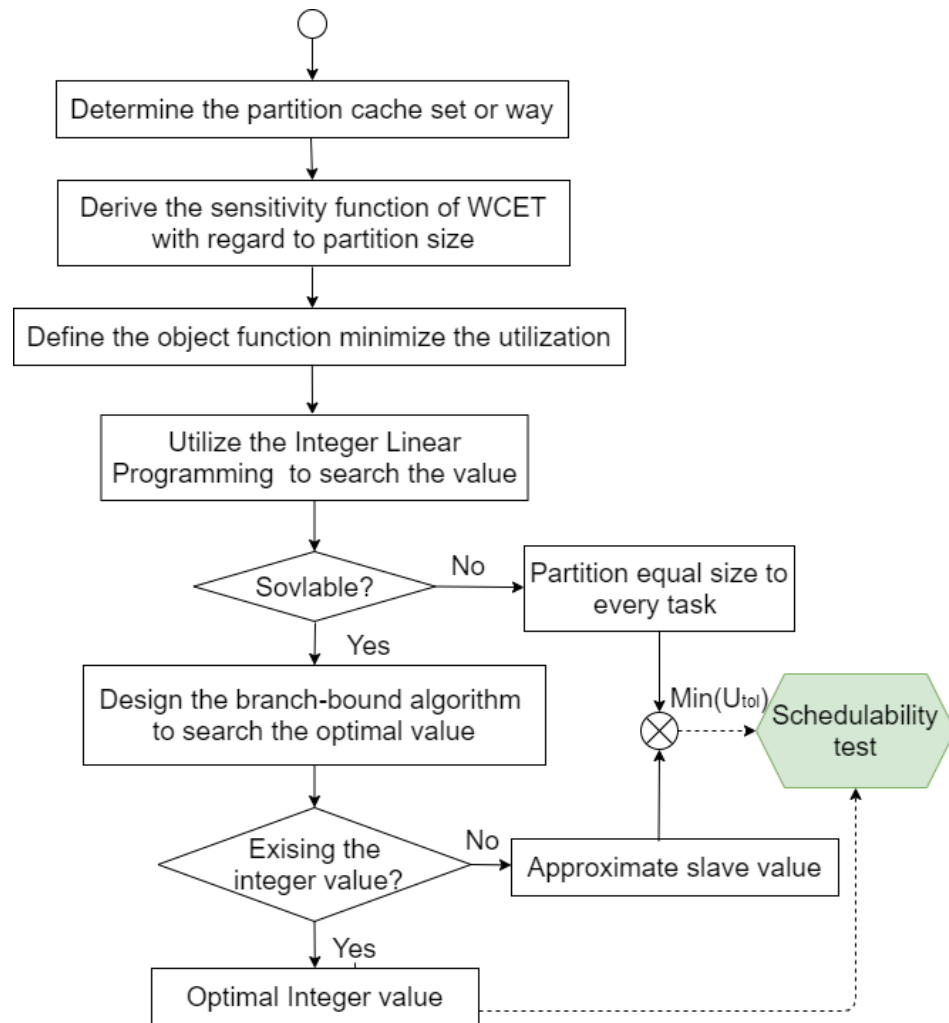


Fig. 6. The comparisons of cache usage and load unbalance for partitioned scheduling policies

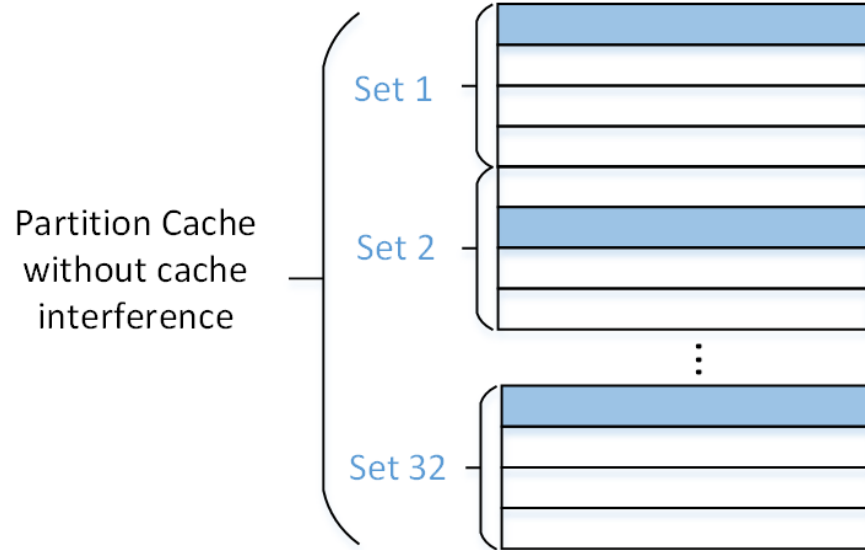
Conclusion

- TCPS: a non-preemptive partitioned scheduling algorithm called TCPS that aims at improving system schedulability through awareness of the cache sensitivity characteristics of applications.
- We have compared the schedulability performance of TCPS with a large range of state-of-the-art approaches, covering all possible scheduling/cache partitioning combinations.


The procedure of cache partition



Partitioning cache way or cache set?



Partition Cache without cache interference



Experiment configuration

Execution platform: an embedded ARM processor with 4 cores.


Cache capacity: L1 private cache--cache line size 16B--512B

L2 shared cache--cache line size 32B--4KB

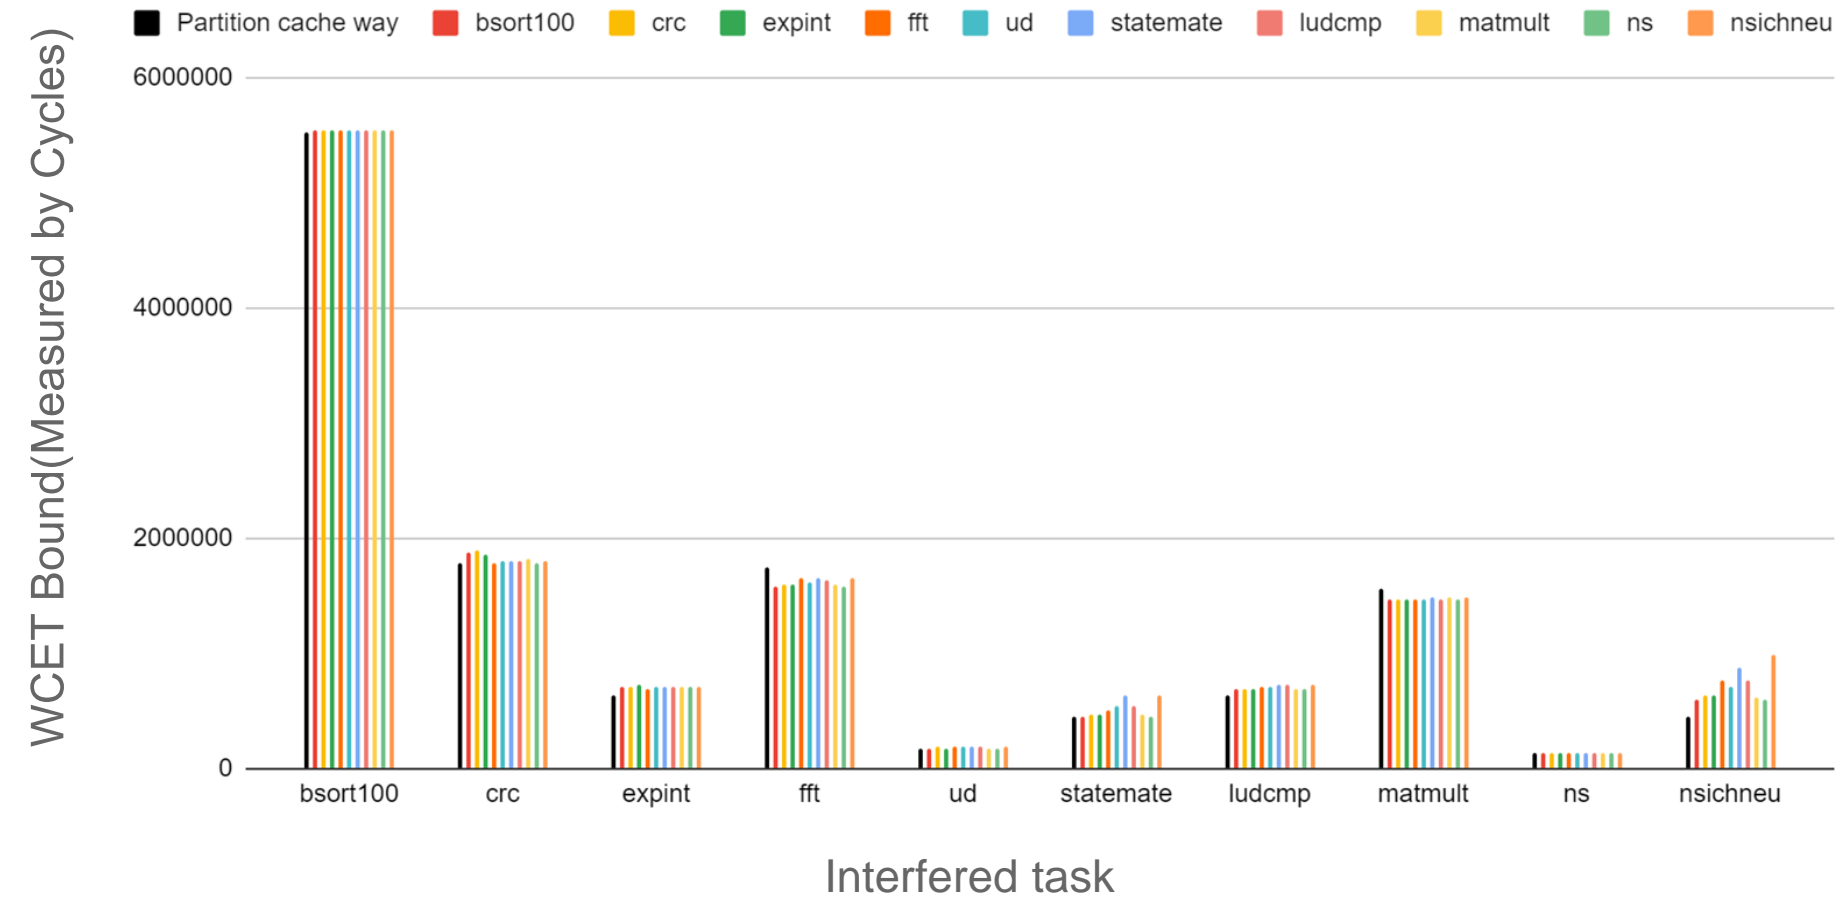
Benchmarks: *Mälardalen* WCET benchmarks

Name	Description
bsort100	bubblesort 100 Integer numbers
crc	Cyclic redundancy check computation
expint	Series expansion for computing an exponential
fft	Fast Fourier Transform
jfdctint	Fast Discrete Cosine Transform
ludcmp	Simultaneous Linear Equations by LU Decomposition
ns	Search in a multi-dimensional array
statemate	Automatically generated code by STARC
matmult	Product of two 20x20 integer matrixes
nsichneu	Simulate an extended Petri net

Partitioning at the granularity of cache set



Partitioning at the granularity of cache way



Experiment configuration

Experiment testbed: an embedded ARM processor with 4 cores.

Cache mode: Direct-mapped Cache

Cache capacity: L1 private cache--cache line size 16B--512B

L2 shared cache--cache line size **32B--8KB**

Benchmarks: **Mälardalen** WCET benchmarks

Name	Description
bsort100	bubblesort 100 Integer numbers
crc	Cyclic redundancy check computation
expint	Series expansion for computing an exponential
fft	Fast Fourier Transform
jfdctint	Fast Discrete Cosine Transform
ludcmp	Simultaneous Linear Equations by LU Decomposition
ns	Search in a multi-dimensional array
statemate	Automatically generated code by STARC
matmult	Product of two 20x20 integer matrixes
nsichneu	Simulate an extended Petri net

The cache partitioning algorithm

$$\min\{U_1 + U_2 + \dots + U_n\}$$
$$\min\left\{\frac{1}{T_1} * c_1 + \frac{1}{T_2} * c_2, \dots, + \frac{1}{T_n} * c_n\right\}$$

Constraints:

$$p_1 + p_2 + \dots + p_n \leq n$$

$$c_1 * A_1 + p_1 * B_1 + k_1 = 0$$

$$c_2 * A_2 + p_2 * B_2 + k_2 = 0$$

...

$$c_n * A_n + p_n * B_n + k_n = 0$$

