PARALLEL COMPUTING SYSTEMS

UNIVERSITEIT VAN AMSTERDAM

Thermal management for S-NUCA Many-Cores via Synchronous Thread Rotations

Presenter: Yixian Shen
Cooperator: Sobhan Niknam
Co-supervisor : Anuj Pathania
Supervisor: Andy D Pimentel

Outline

- ***Research background***
- A motivational example
- The key designs of our approach
- Experimental results
- Conclusion

Research background→S-NUCA architecture

- Static Non-uniform Cache Architecture

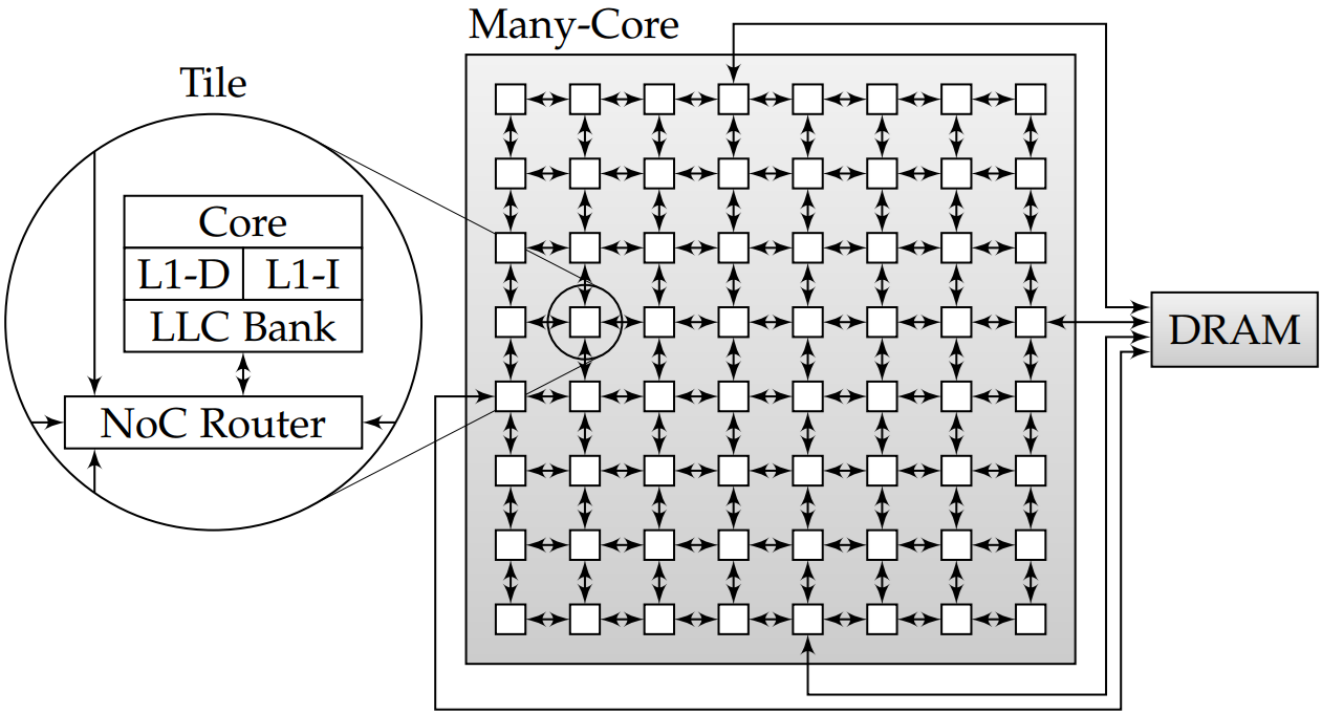

M. Rapp, A. Pathania, T. Mitra and J. Henkel, "**Neural Network-Based Performance Prediction for Task Migration on S-NUCA Many-Cores**," in *IEEE Transactions on Computers*,

Research background→S-NUCA architecture

- Characteristics of S-NUCA
  - Logically-shared yet physically-distributed.

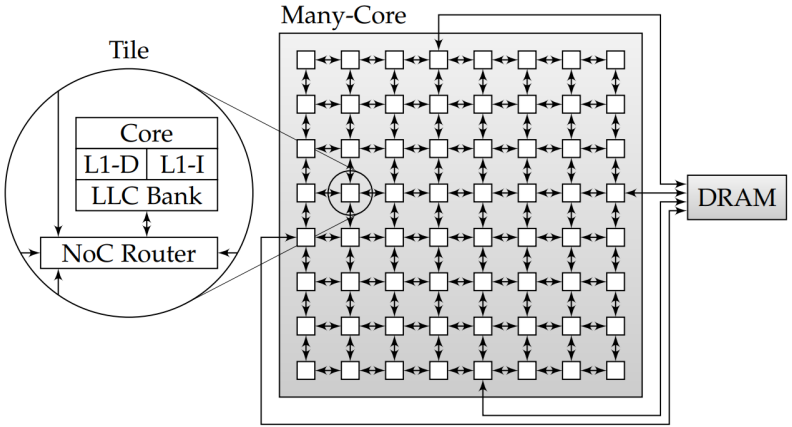

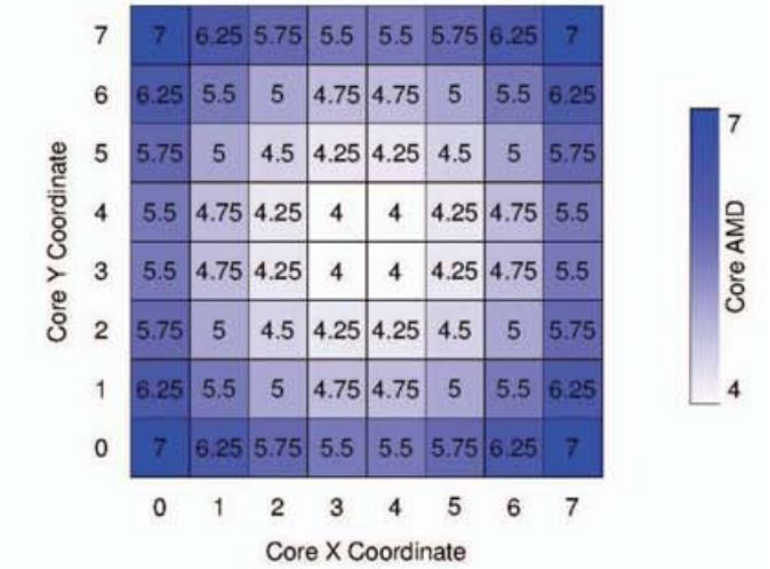

A. Pathania and J. Henkel, "**Task scheduling for many-cores with S-NUCA caches**," *2018 Design, Automation & Test in Europe Conference & Exhibition (DATE)*, 2018, pp. 557-562, doi: 10.23919/DATE.2018.8342069

Dark Silicon in S-NUCA many-cores

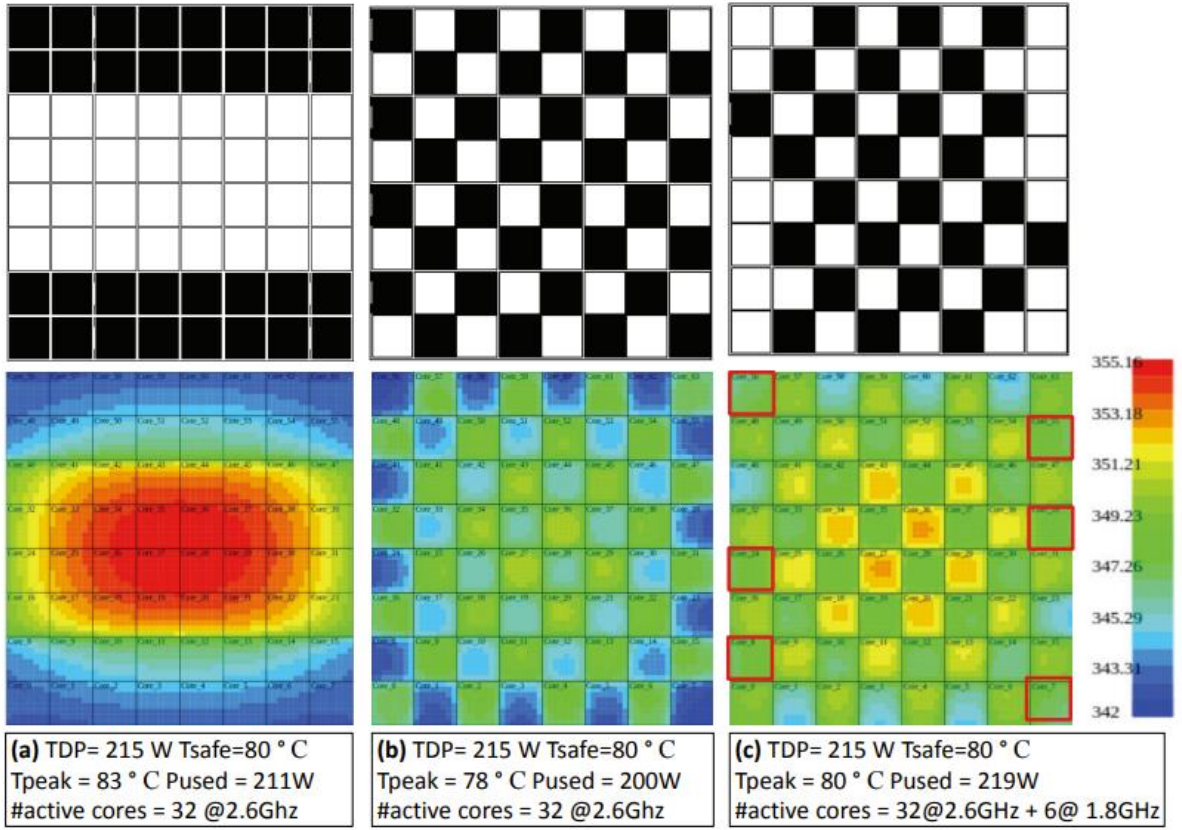

Shafique, Muhammad, et al. "**Dark silicon as a challenge for hardware/software co-design:** Invited special session paper." Proceedings of the 2014 International Conference on Hardware/Software Codesign and System Synthesis. 2014..

Research background →R-C thermal model

- Thermal behavior is dual to electronic circuit.

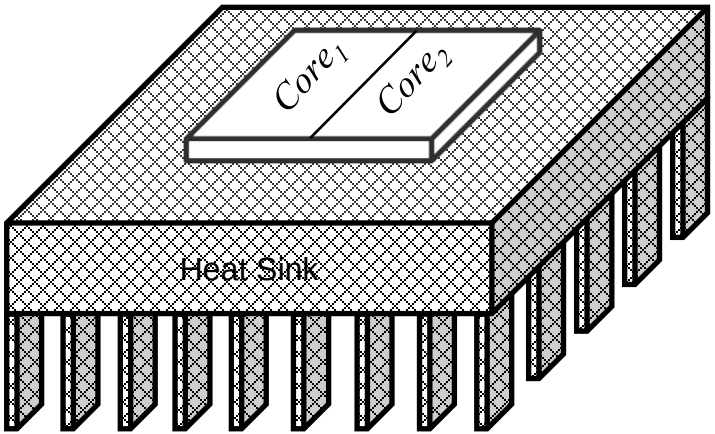

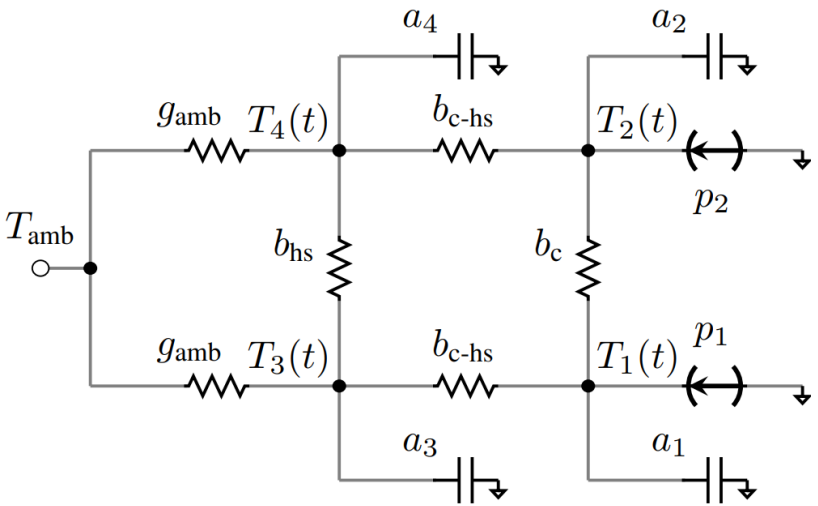

Research background →R-C thermal model

- Thermal behavior is dual to electronic circuit.

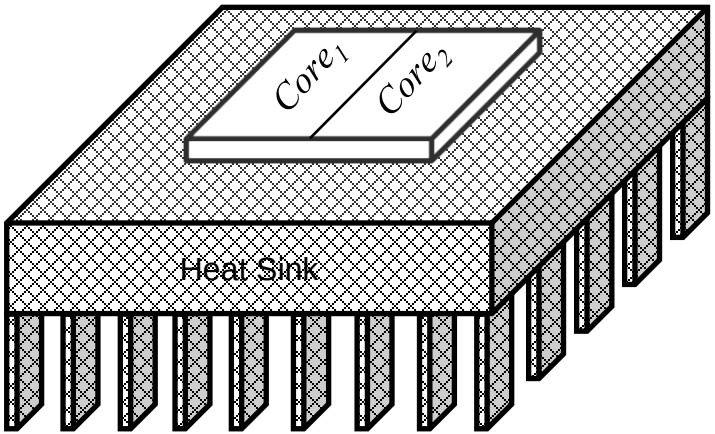

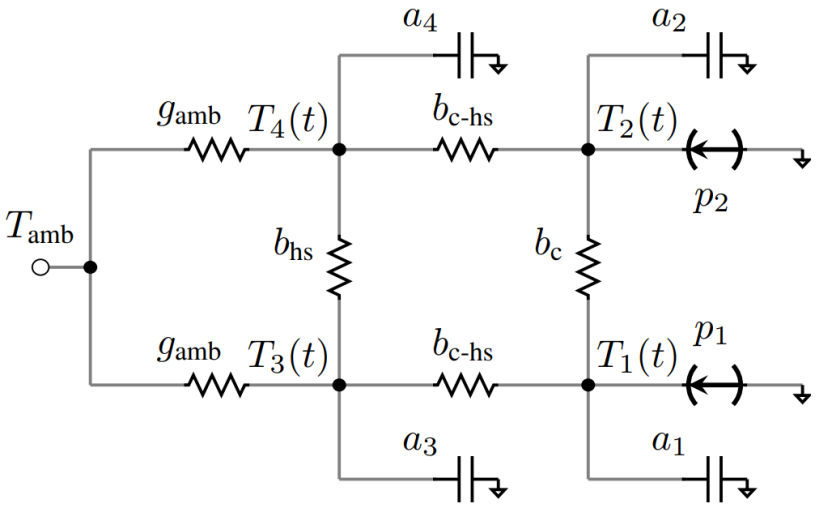

$$\mathbf{A}\mathbf{T}' + \mathbf{B}\mathbf{T} = \mathbf{P} + T_{\text{amb}}\mathbf{G}$$

Outline

- Research background
- ***A motivational example***
- The key designs of our approach
- Experimental results
- Conclusion.

A motivational example

- Experimental platform: simulated 16-core S-NUCA processors
- Experimental configuration
- Thermal threshold 70(℃)

	Benchmark	Thread	DTM method	Execute at
Case 1	blackscholes	1 master, 1 slave	No	Core 5,10
Case 2	blackscholes	1 master, 1 slave	TSP	Core 5,10
Case 3	blackscholes	1 master, 1 slave	Task rotation	Core 5,6,9,10

A motivational example

	Benchmark	Thread	DTM method	Execute at
Case 1	blackscholes	1 master, 1 slave	No	Core 5,10
Case 2	blackscholes	1 master, 1 slave	TSP	Core 5,10
Case 3	blackscholes	1 master, 1 slave	Task rotation	Core 5,6,9,10

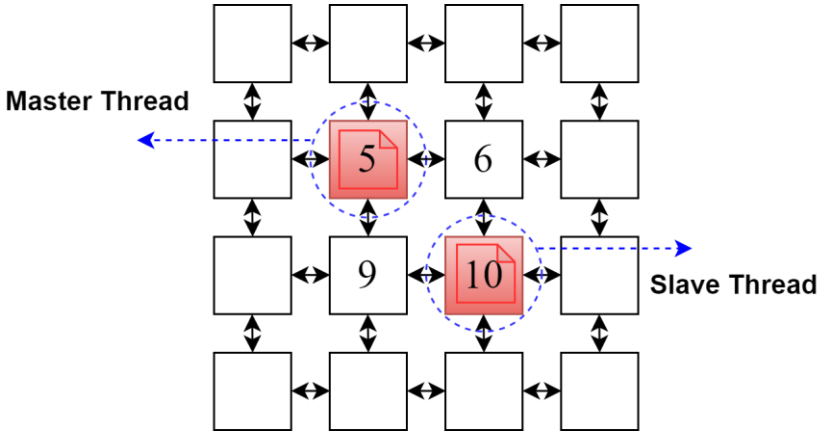

Case1 and Case 2

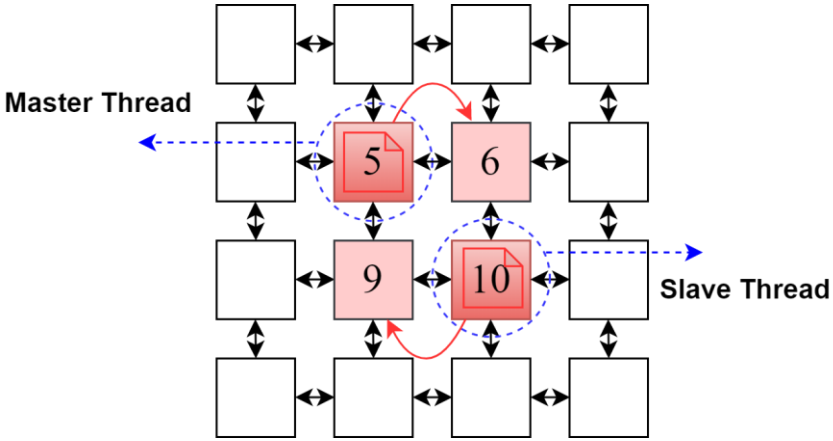

Case 3

A motivational example---Case 1

- Peak temperature 80.03(℃)
- Execution time 67.97(ms)

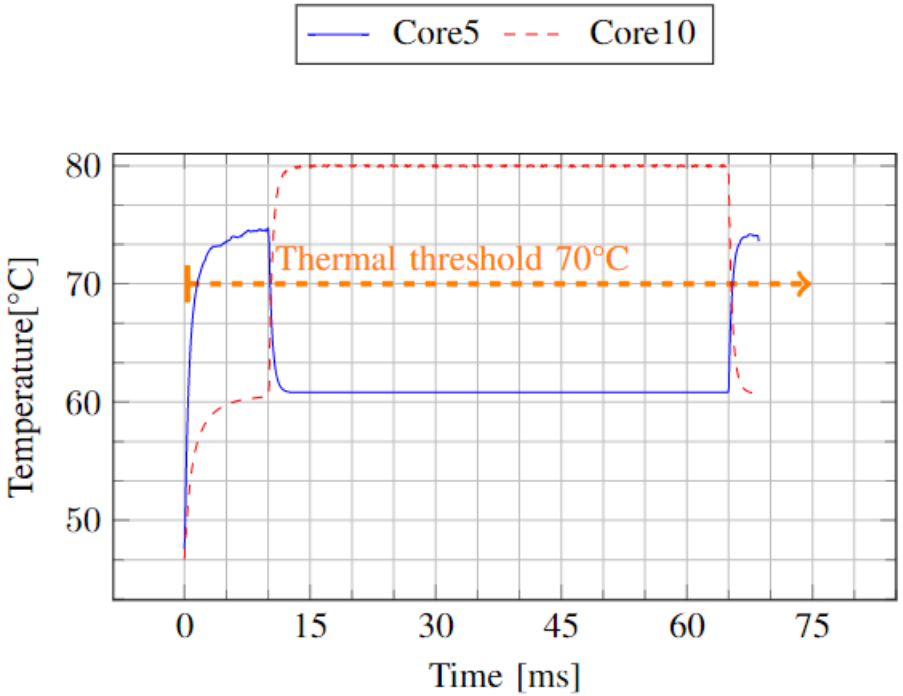

A motivational example---Case 2

- Peak temperature 67.94(℃)
- Execution time 84.49(ms)

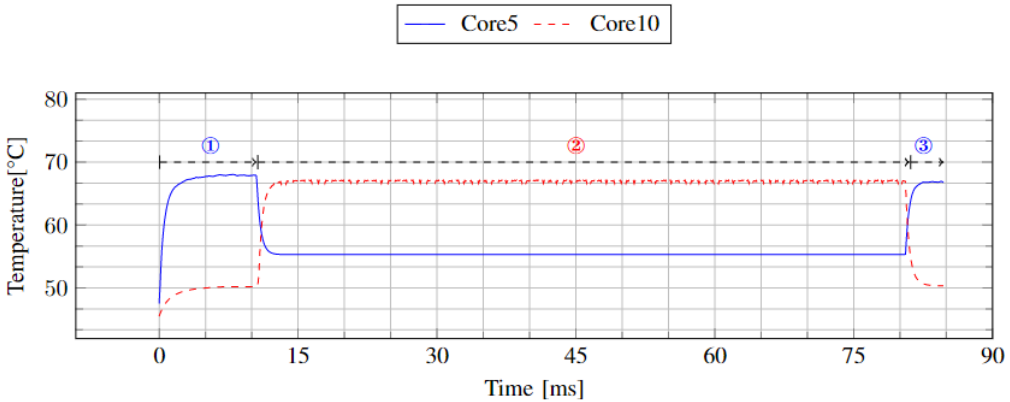

(a) Thermally Safe Frequency provided by *TSP* [11]

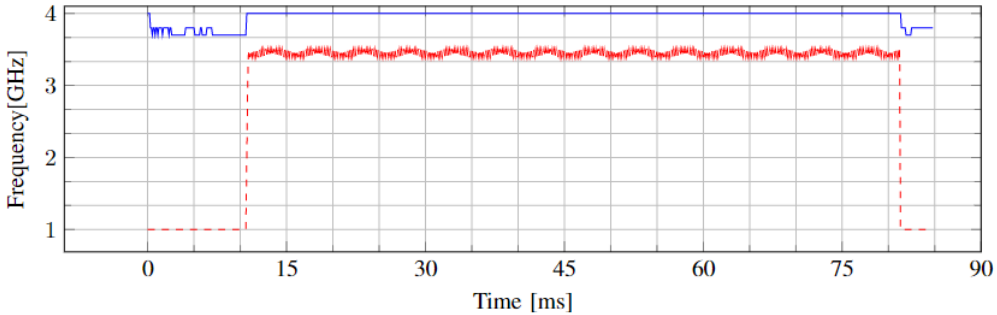

(b) Frequency profiling by TSP

A motivational example---Case 3

- Peak temperature 69.32(℃)
- Execution time 74.47(ms)

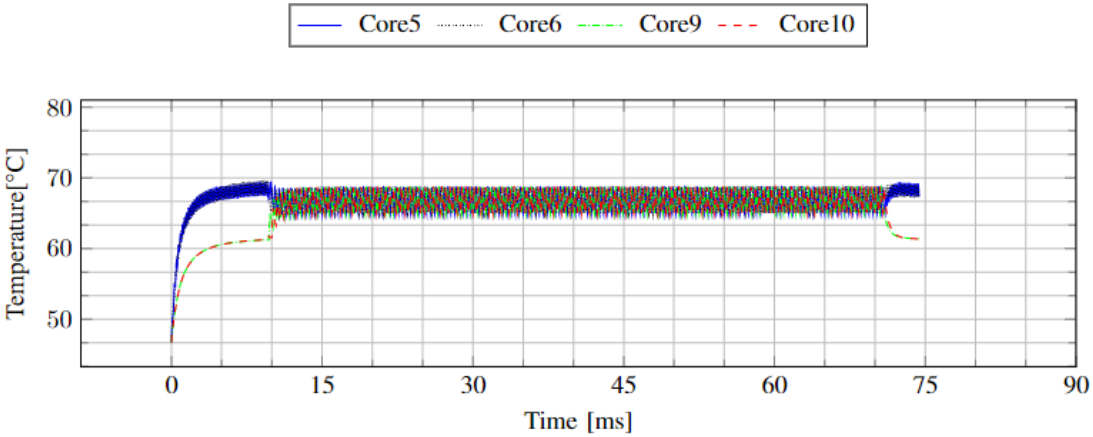

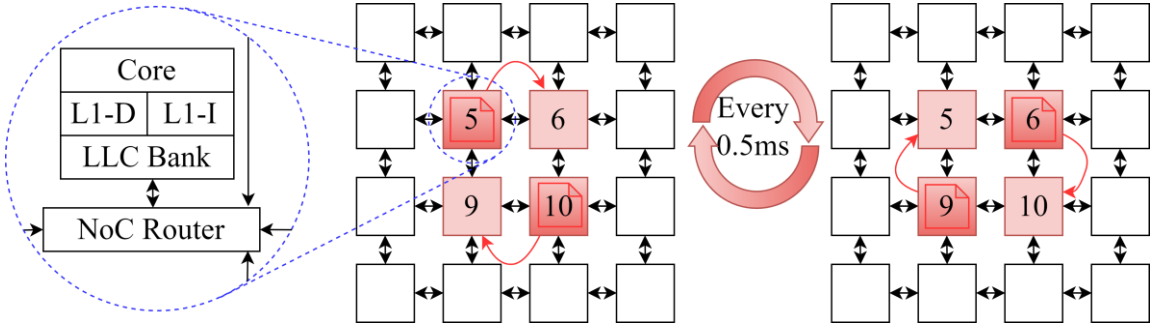

A motivational example

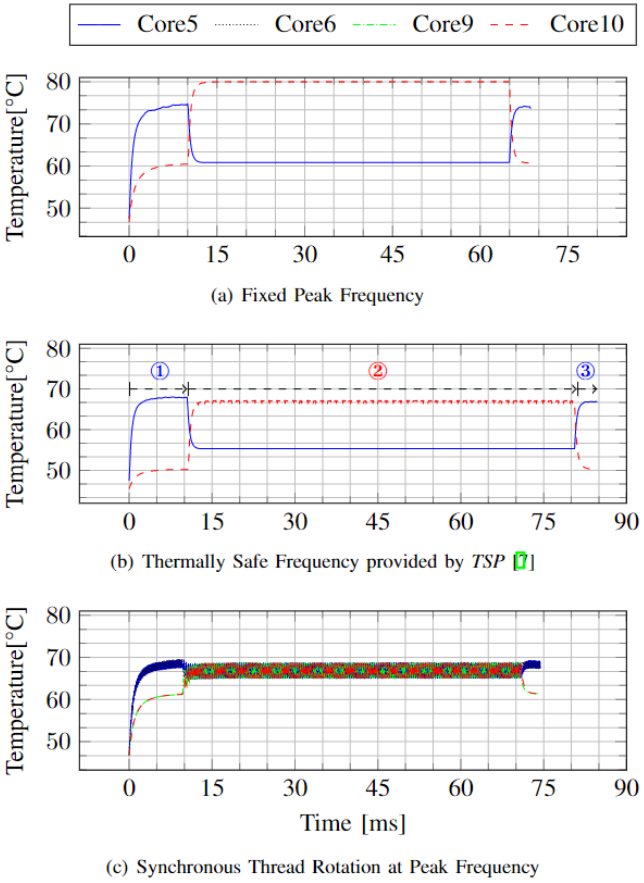

(a) Fixed Peak Frequency

(b) Thermally Safe Frequency provided by *TSP* [7]

(c) Synchronous Thread Rotation at Peak Frequency

	Benchmark	Thread	DTM method	Execute at
Case 1	blackscholes	1 master, 1 slave	No	Core 5,10
Case 2	blackscholes	1 master, 1 slave	TSP	Core 5,10
Case 3	blackscholes	1 master, 1 slave	Task rotation	Core 5,6,9,10

	Case 1	Case 2	Case 3
Peak temp(℃)	80.03	67.94	69.32
Exec time(ms)	67.97	84.49	74.47
Penalty(%)		19.55	8.72

Case 3 is **10.83% faster** than Case 2

Rotation penalty **<** DVFS-based penalty

Outline

- Research background
- A motivational example
- ***The key designs of our approach***
- Experimental results
- Conclusion.

The key designs of our approach

- Thermal regulations
  - A theoretical one-shot peak temperature calculation
- Architecture-aware task scheduling
  - *Hotpotato* thermal-aware scheduler.

The key designs of our approach

- Thermal regulations
  - A theoretical one-shot peak temperature calculation

$$\mathbf{T}_{(d\delta+1)\tau} = \mathbf{V}\cdot\mathbf{diag}\left(\frac{1}{1-e^{\delta\lambda_1\tau}},\frac{1}{1-e^{\delta\lambda_2\tau}},\cdots,\frac{1}{1-e^{\delta\lambda_N\tau}}\right)\cdot\mathbf{V}^{-1}\mathbf{wP}_{\tau}$$
$$+\mathbf{V}\cdot\mathbf{diag}\left(\frac{e^{(\delta-1)\lambda_1\tau}}{1-e^{\delta\lambda_1\tau}},\frac{e^{(\delta-1)\lambda_2\tau}}{1-e^{\delta\lambda_2\tau}},\cdots,\frac{e^{(\delta-1)\lambda_N\tau}}{1-e^{\delta\lambda_N\tau}}\right)\cdot\mathbf{V}^{-1}\mathbf{wP}_{2\tau}+\cdots$$
$$+\mathbf{V}\cdot\mathbf{diag}\left(\frac{e^{\lambda_1\tau}}{1-e^{\delta\lambda_1\tau}},\frac{e^{\lambda_2\tau}}{1-e^{\delta\lambda_2\tau}},\cdots,\frac{e^{\lambda_N\tau}}{1-e^{\delta\lambda_N\tau}}\right)\cdot\mathbf{V}^{-1}\mathbf{wP}_{\delta\tau} \tag{10}$$

The key designs of our approach

- Thermal regulations
  - A theoretical one-shot peak temperature calculation

$$\mathbf{T}_{(d\delta+1)\tau} = \mathbf{V}\cdot\mathbf{diag}\left(\frac{1}{1-e^{\delta\lambda_1\tau}},\frac{1}{1-e^{\delta\lambda_2\tau}},\cdots,\frac{1}{1-e^{\delta\lambda_N\tau}}\right)\cdot\mathbf{V}^{-1}\mathbf{w}\mathbf{P}_{\tau}$$
$$+\mathbf{V}\cdot\mathbf{diag}\left(\frac{e^{(\delta-1)\lambda_1\tau}}{1-e^{\delta\lambda_1\tau}},\frac{e^{(\delta-1)\lambda_2\tau}}{1-e^{\delta\lambda_2\tau}},\cdots,\frac{e^{(\delta-1)\lambda_N\tau}}{1-e^{\delta\lambda_N\tau}}\right)\cdot\mathbf{V}^{-1}\mathbf{w}\mathbf{P}_{2\tau}+\cdots$$
$$+\mathbf{V}\cdot\mathbf{diag}\left(\frac{e^{\lambda_1\tau}}{1-e^{\delta\lambda_1\tau}},\frac{e^{\lambda_2\tau}}{1-e^{\delta\lambda_2\tau}},\cdots,\frac{e^{\lambda_N\tau}}{1-e^{\delta\lambda_N\tau}}\right)\cdot\mathbf{V}^{-1}\mathbf{w}\mathbf{P}_{\delta\tau} \quad (10)$$

Auxiliary matrix using floorplan-based constants

Rotation speed

Rotation period

Power traces

The key designs of our approach

- Thermal regulations
  - A theoretical one-shot peak temperature calculation

$$T_{peak} = f(v, w, \delta, \tau, P)$$

Auxiliary matrix using floorplan-based constants

Rotation speed

Rotation period

Power traces

The hotpotato scheduling

- The inherent heterogeneity of S-NUCA

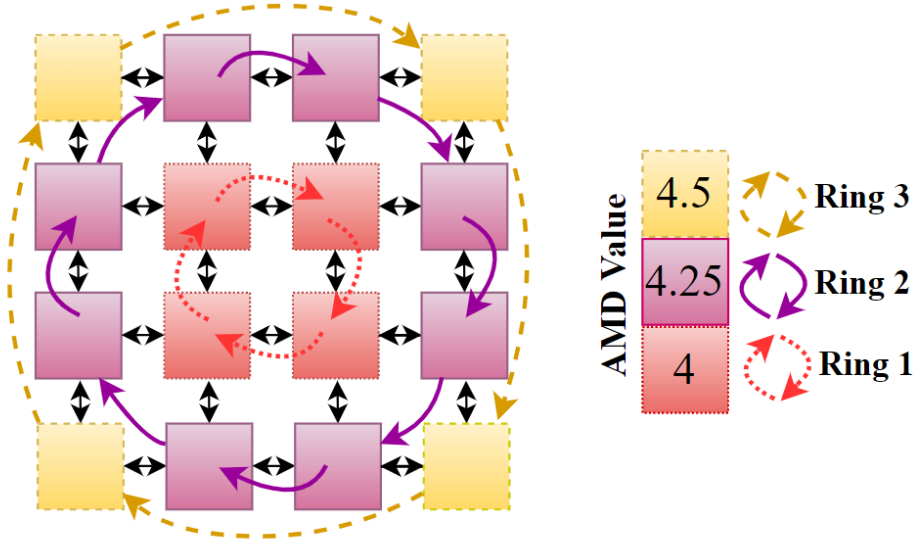

Fig. 3: Abstraction for concentric AMD-based rotation rings.

- Access cache latency
  Ring1 < Ring2 < Ring3
- Thermal dissipation condition
  Ring1 < Ring2 < Ring3

The hotpotato scheduling

- ***maximizing the performance*** under thermal threshold
  - Keep peak temperature T_{peak} close to thermal threshold T_{DTM}
    - Migrate to the thread to lower ring close to the center ring
  - Minimize the task rotation penalty
    - Slow down the rotation speed

The hotpotato scheduling

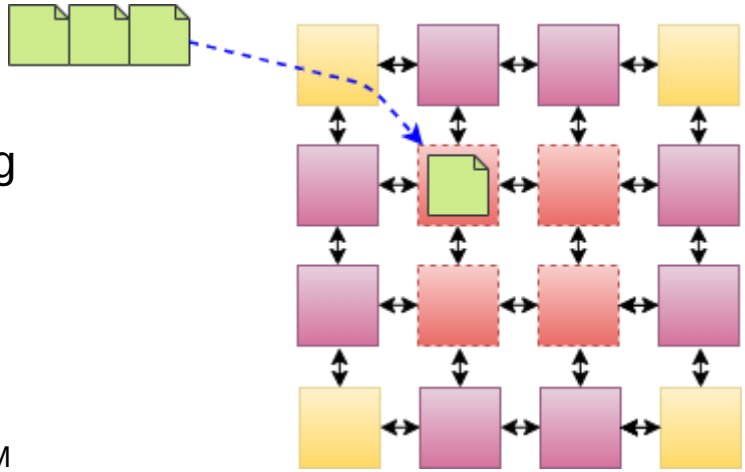

- New threads enter S-NUCA
  - ❶ Attempt to assign thread to the lowest ring
  - ❷ Calculate the T_{peak}
  - ❸ If $T_{peak} < T_{DTM}$, rotate at that ring
  - ❹ Else, attempt to L+1 rings until $T_{peak} < T_{DTM}$
  - ❺ Reach R_{max} ,still not managed, speed up the rotation speed

The hotpotato scheduling

- Thermal violation during runtime

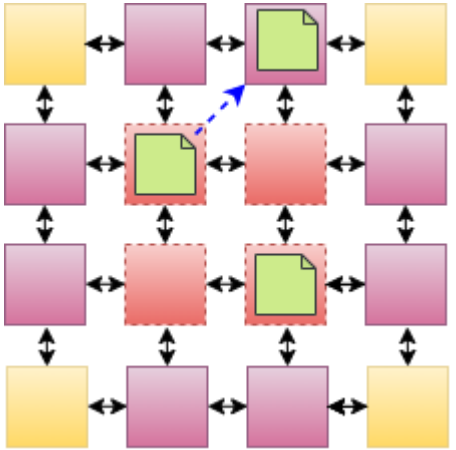

❶ Rank the threads by CPI(cycle per instruction)

❷ Select the thread with lowest CPI

❸ Migrate thread to the L+1 rings

❹ If $T_{peak} < T_{DTM}$, rotate at that ring

❺ Else, attempt to L+1 rings until $T_{peak} < T_{DTM}$

❻ Reach R_{max} ,still not managed, speed up the rotation speed

The hotpotato scheduling

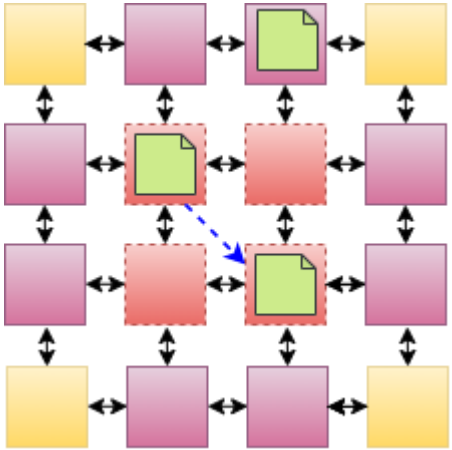

- Headroom not exploited
  - ❶ Rank the threads by CPI(cycle per instruction)
  - ❷ Select the thread with highest CPI
  - ❸ Migrate thread to the L-1 rings
  - ❹ If $T_{DTM} - T_{peak} < \Delta$, rotate at that ring
  - ❺ Else, attempt to L-1 rings until $T_{DTM} - T_{peak} < \Delta$
  - ❻ Reach R_{min} ,still not managed, slow down the rotation speed

Outline

- Research background
- A motivational example
- The key designs of our approach
- ***Experimental results***
- Conclusion.

The experimental configuration

Core Parameter	
Number of Cores	64
Core Model	x86, 4.0 GHz, 14 nm, out-of-order
L1 I/D cache	16/16 KB, 8/8-way,64B-block
LLC	128 KB per core, 16-way, 64B-block
NoC Latency	1.5ns per hop
Noc link width	256 Bit
The area of core	0.81 mm^2

The experimental results

- Full loaded system

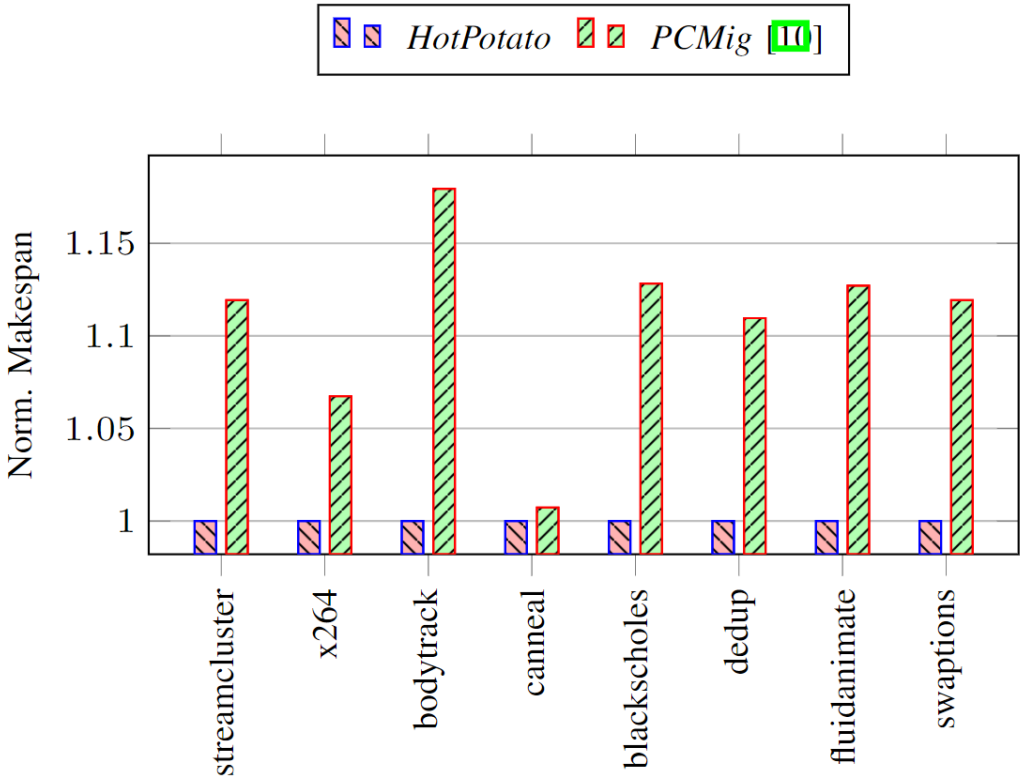

(a) Homogeneous Workload

The experimental results

- Open system
- The arrival time follow Poisson distribution.

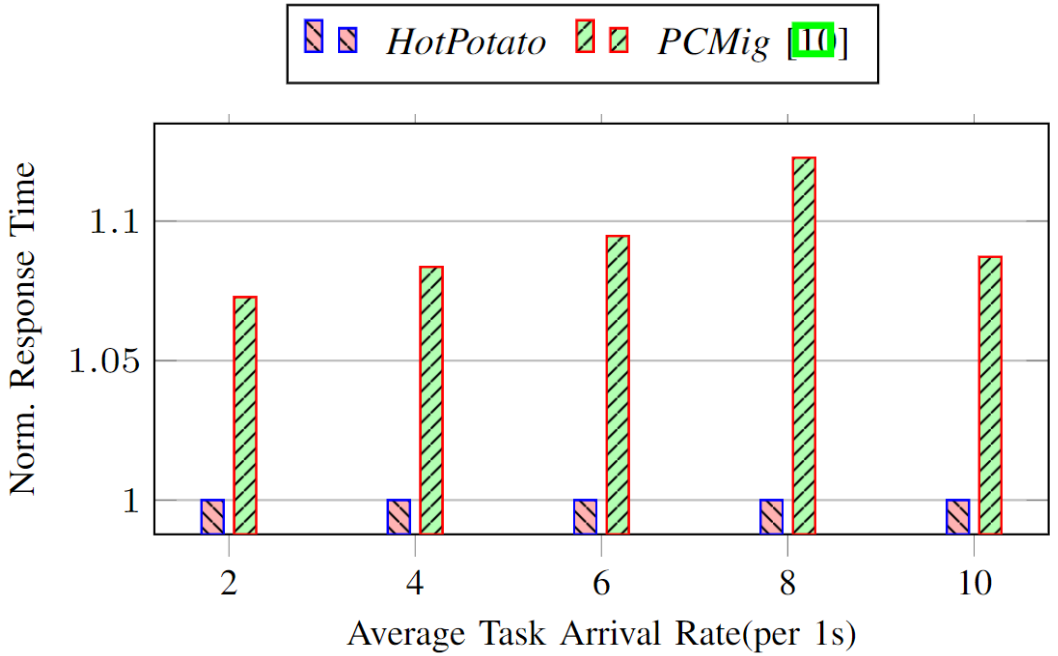

(b) Heterogeneous Workload

Outline

- Research background
- A motivational example
- The key designs of our approach
- Experimental results
- ***Conclusion.***

The conclusion.

- We are the first to use synchronous thread rotations for thermal management.
- We propose a one-shot peak temperature calculation approach based on the RC thermal model.
- We present a *hotpotato* scheduling for task scheduling to speed up the performance.
- We implement our hotpotato scheduler in the state-of-the-art thermal simulator hotsniper

Thanks for your attention

y.shen@uva.nl

Questions

The hotpotato scheduling

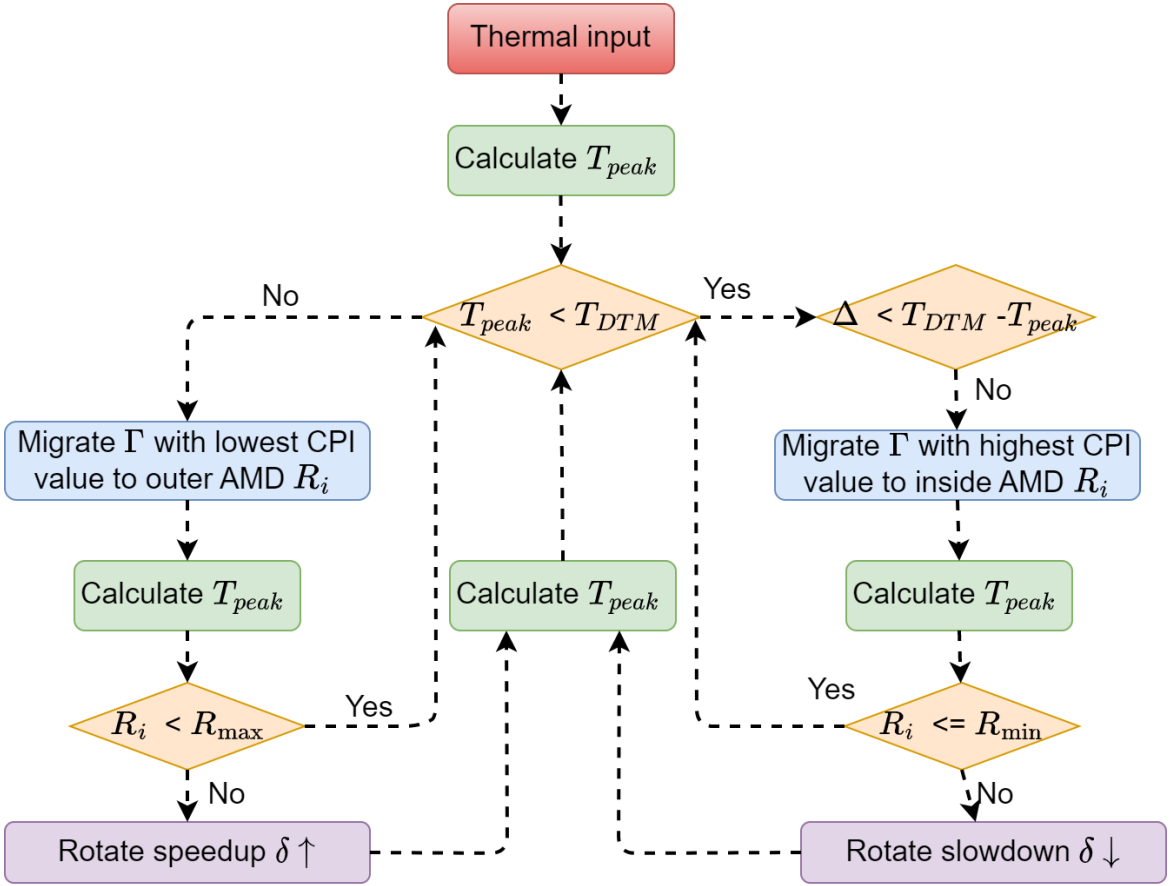

The hotpotato scheduling

- Thermal emergency may happen

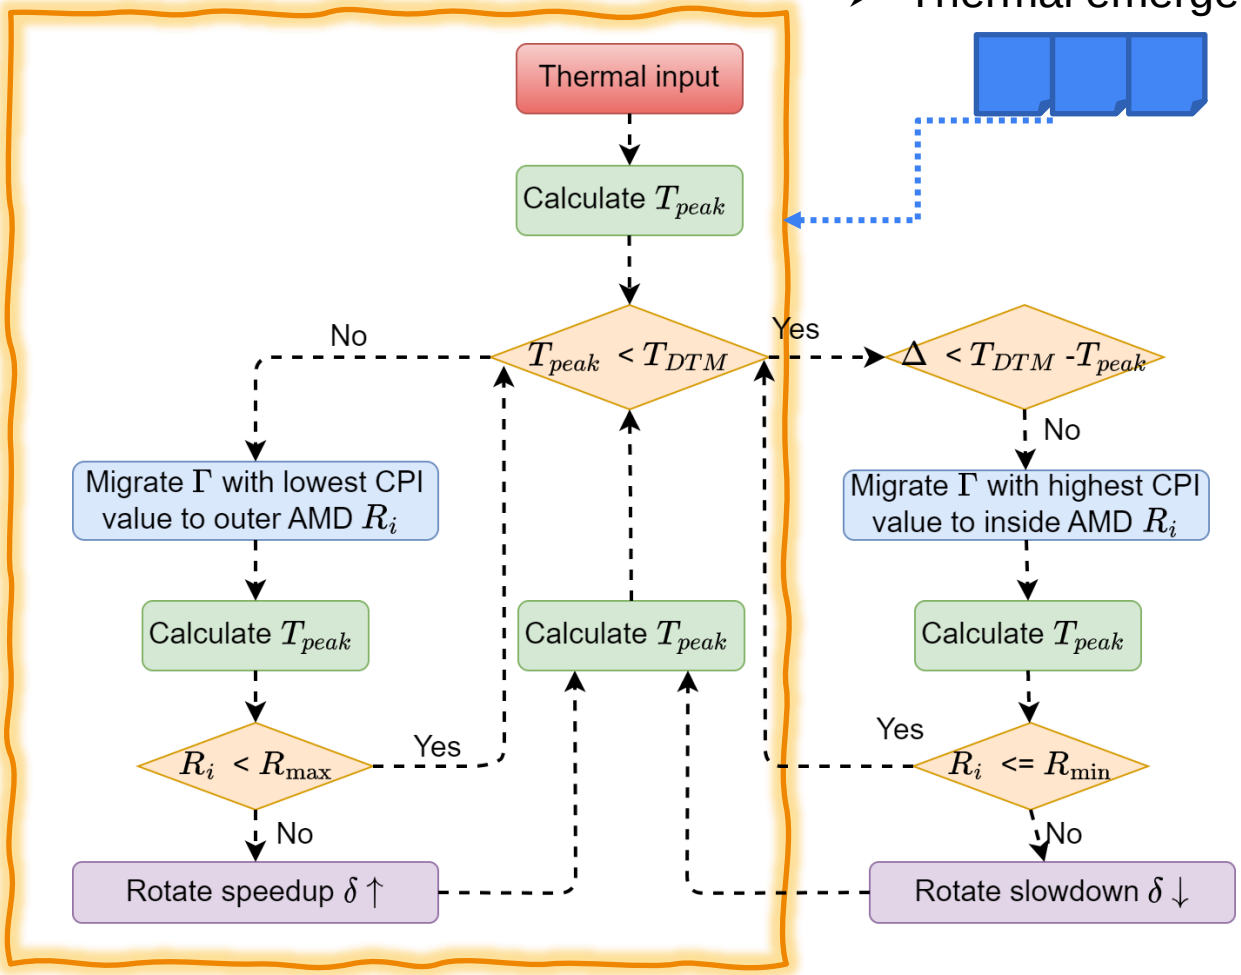

The hotpotato scheduling

- Thermal headroom not exploited

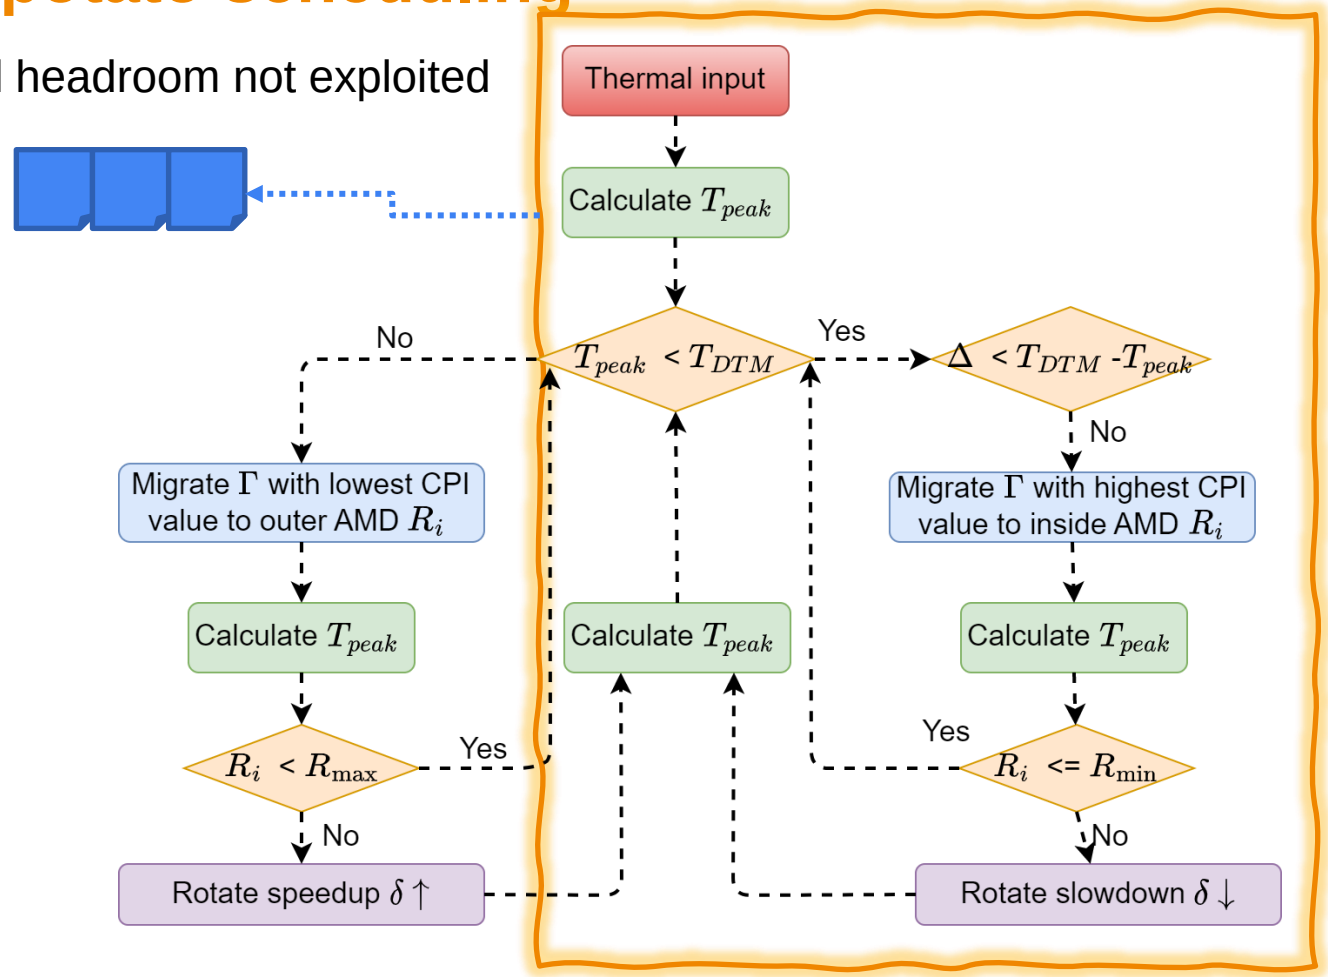

The limitations of this work.

- Only suitable for the system with LLC.
- The heuristic of hotpotato is a little bit straightforward.
- Not sure if the data congestion crashes the system when executing on more core systems.

The hotpotato scheduling

- Bohr's model

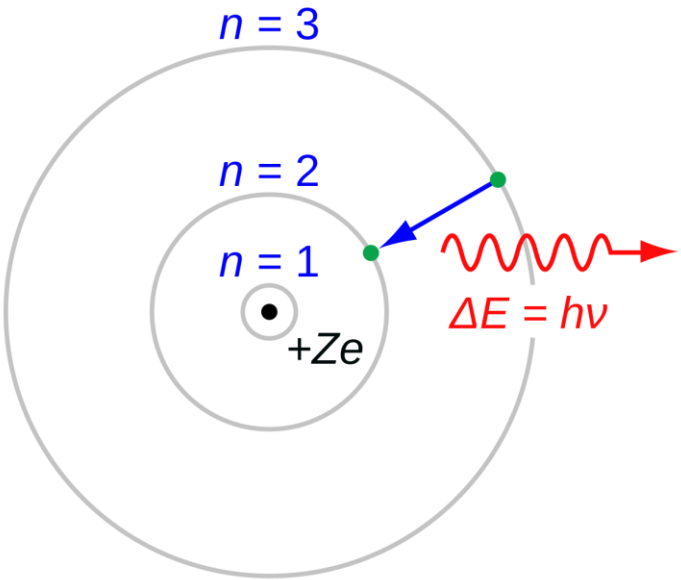

- Every orbit with different energy level
- The atom gains energy
  Jumping from low orbit to high orbit
- The atom lose energy
  falling from high orbit to low orbit